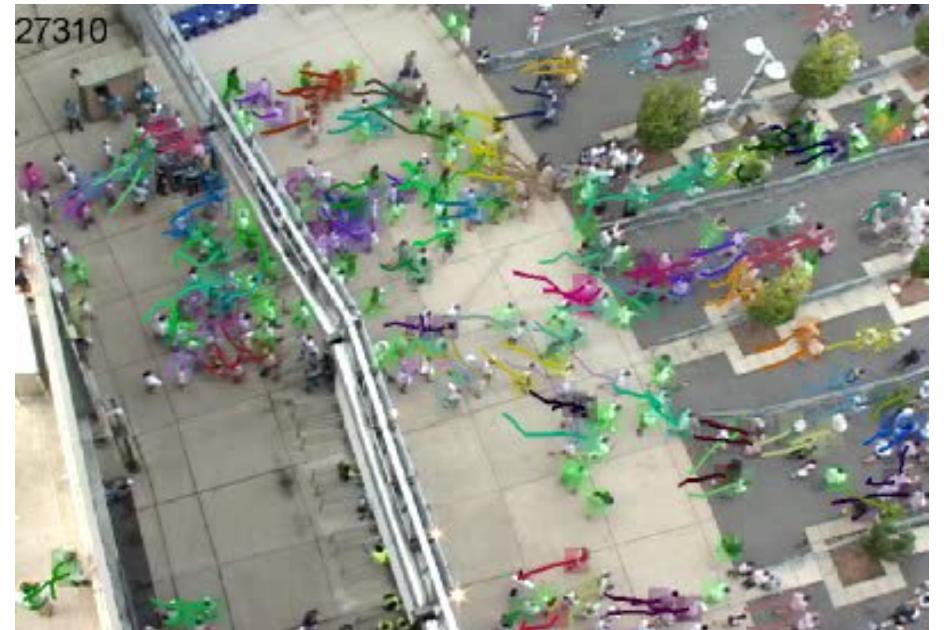


Crowd Scene Analysis

- Using computer vision tools to look at people in public places
- Real-time monitoring
 - situation awareness
 - notifications/alarms
- After-action review
 - traffic analysis



Crowd Scene Analysis

Things we might want to know:

- How many people are there?
- How to track specific individuals?
- How to determine who is with whom?

Challenges:

Crowd scenes tend to have low resolution.

You rarely see individuals in isolation.

Indeed, there are frequent partial occlusions.

Crowd Counting

FAQ: How many people participated in ...

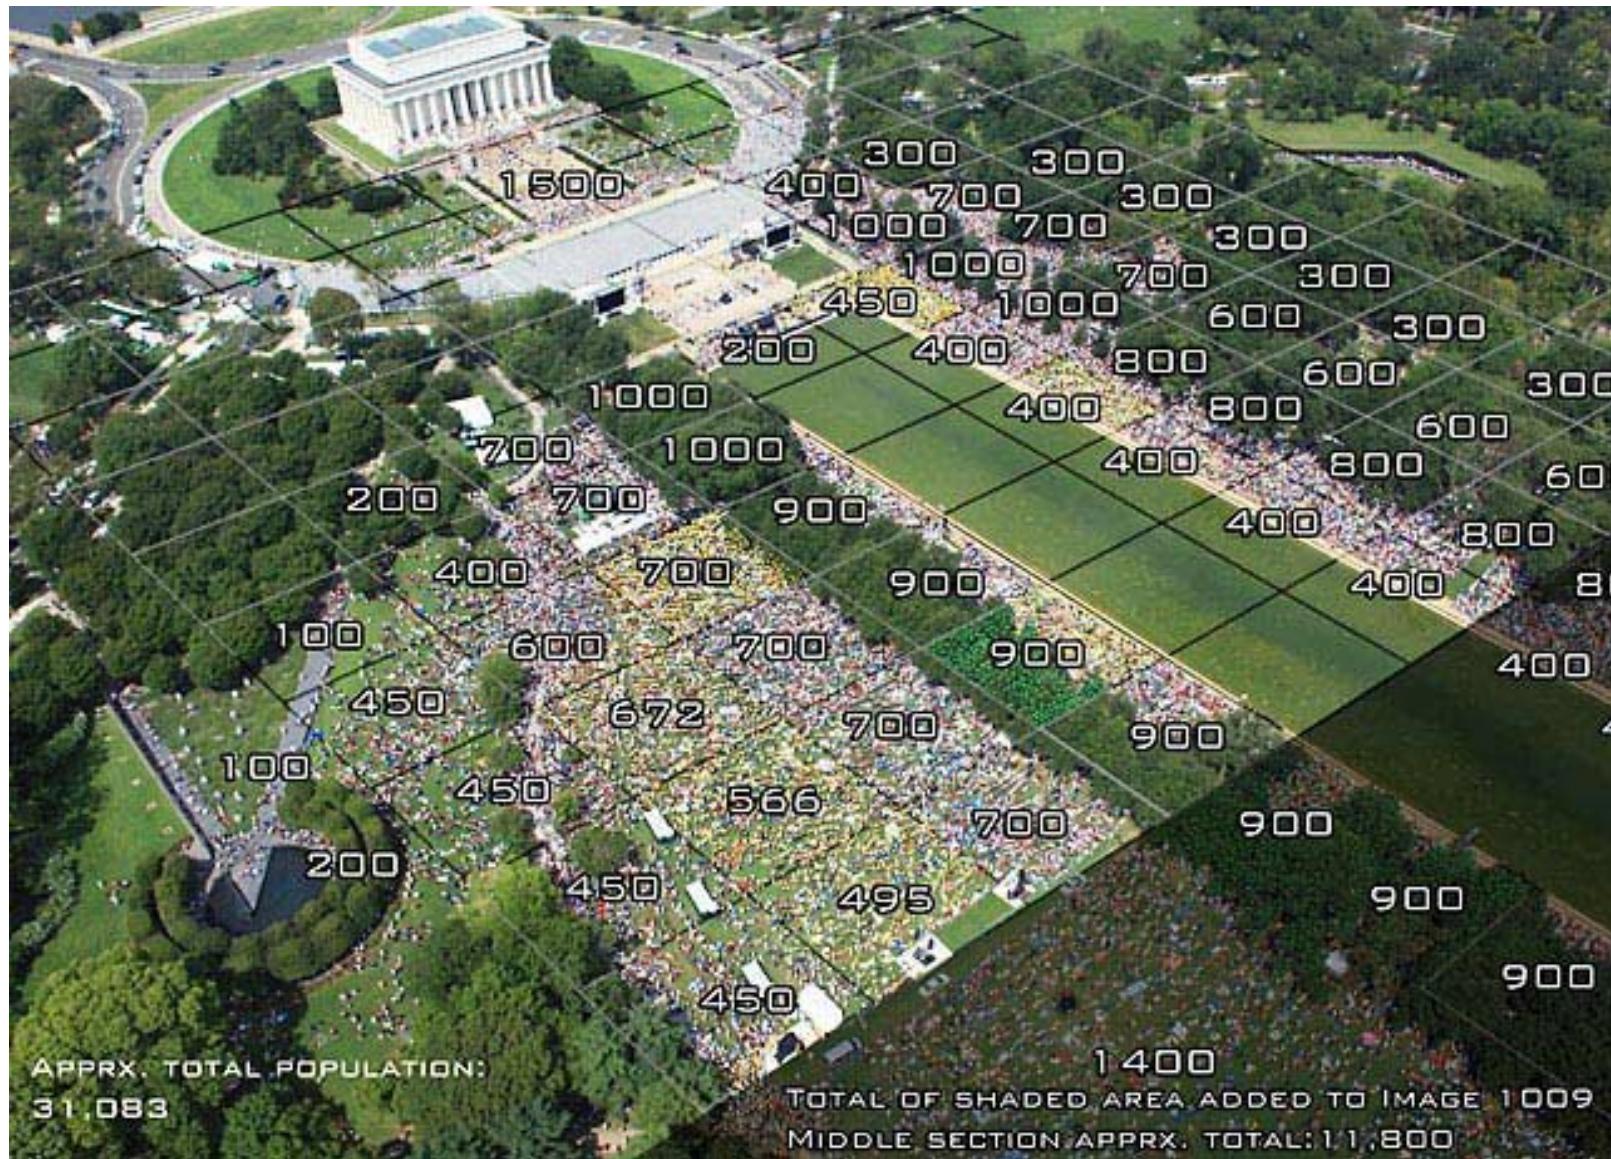
- Tahrir Square Protests
- Obama's inauguration
- Occupy Wall Street
- Kumbh Mela



Jacob's Method

- Herbert Jacobs, Berkeley, 1960s
- $\text{count} = \text{area} * \text{density}$
 - 10 sqft/person – loose crowd (arm's length from each other)
 - 4.5 sqft/person – more dense
 - 2.5 sqft/person – very dense (shoulder-to-shoulder)
- Problem: Pedestrians do not uniformly distribute over a space, but clump together into groups or clusters.
- Refinement: break area into a grid of ground patches and estimate a different density in each small patch. Accumulate these counts over whole area.

Example of Jacob's Method



source <http://www.popularmechanics.com/science/the-curious-science-of-counting-a-crowd>

Computer Vision Could do Better!

Cavaet: nobody really wants accurate counts

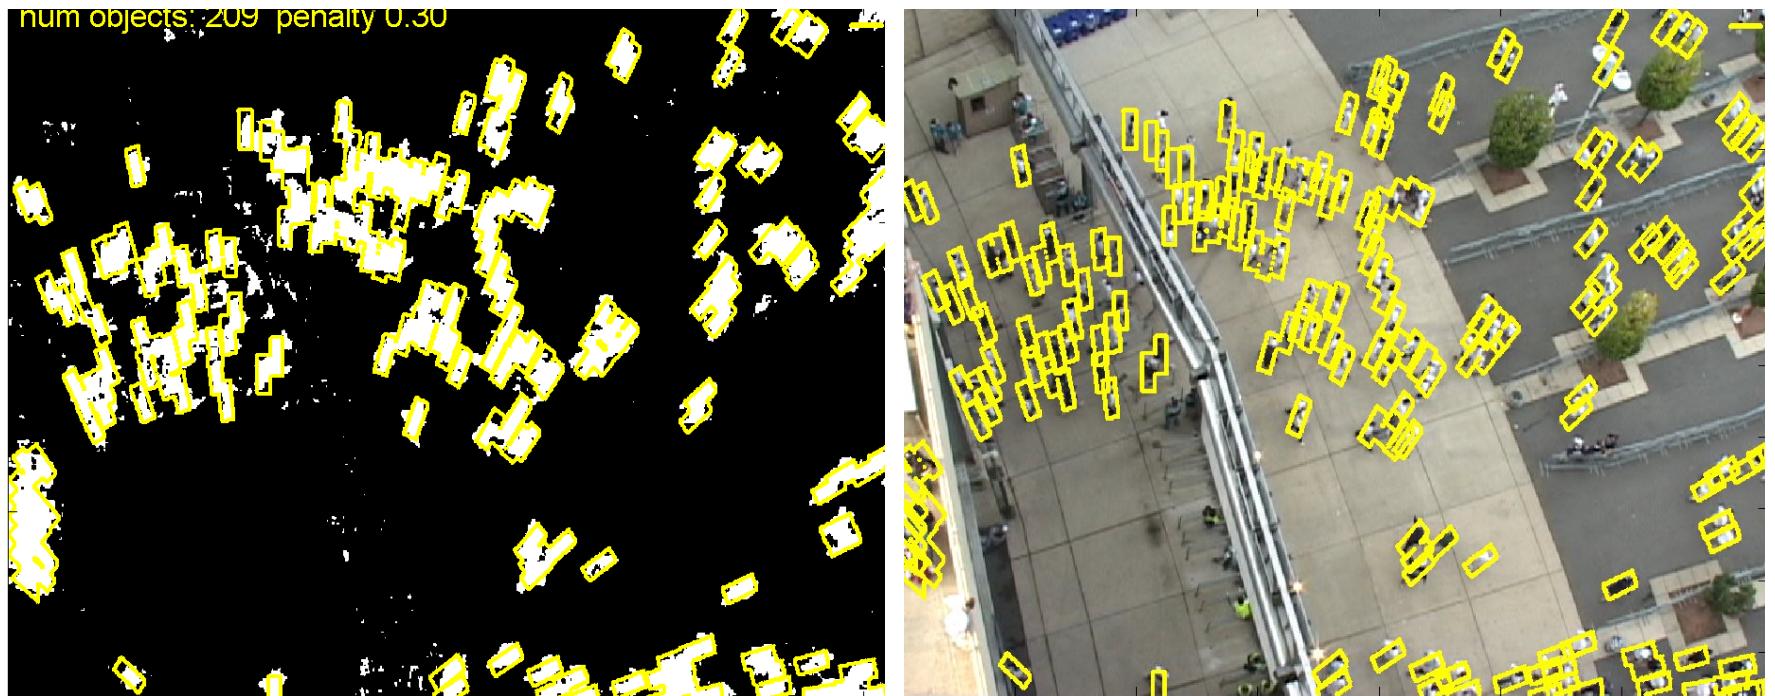
e.g. organizers of the “Million Man March” in Washington DC threatened to sue the National Park Service for estimating that only 400K people attended.

Vision-based Counting

- detection and tracking (light density)
- clustering feature trajectories that move coherently (moderate density)
- treat crowd as a dynamic texture and compute regression estimates based on measured properties (heavy density)

Detecting and Counting Individuals

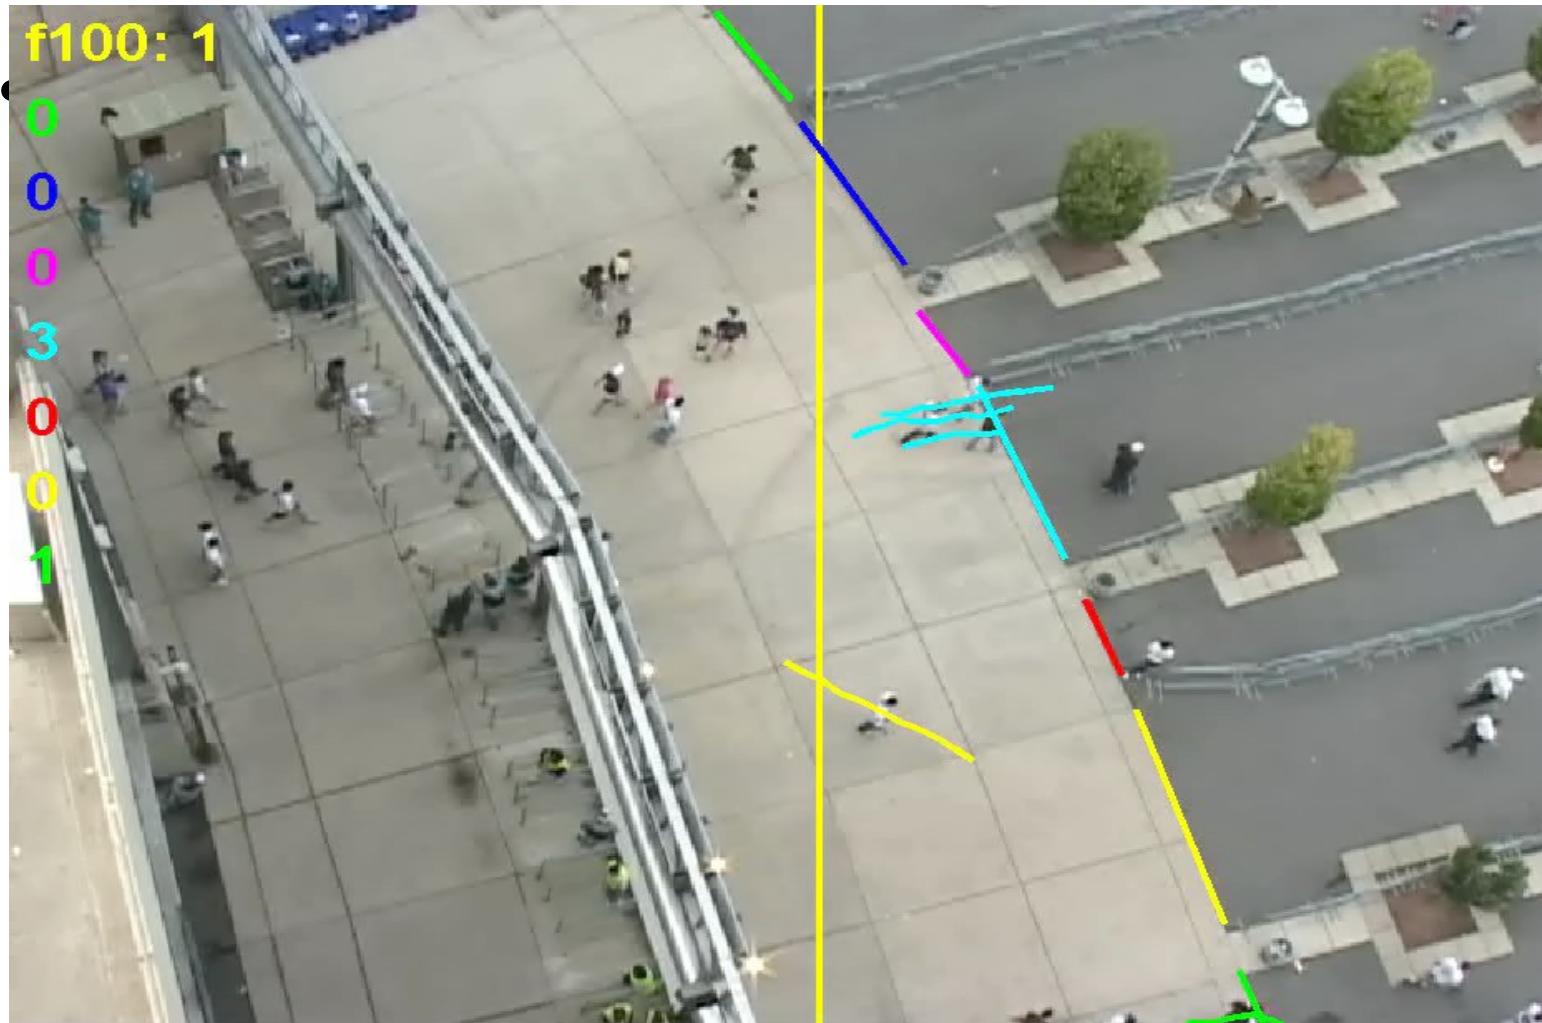
Ge and Collins, "Marked Point Processes for Crowd Counting," *IEEE Computer Vision and Pattern Recognition (CVPR'09)*, Miami, FL, June 2009, pp.2913-2920.



Good for low-resolution / wide-angle views.
Relies on foreground/background segmentation.
Not appropriate for very high crowd density or stationary people.

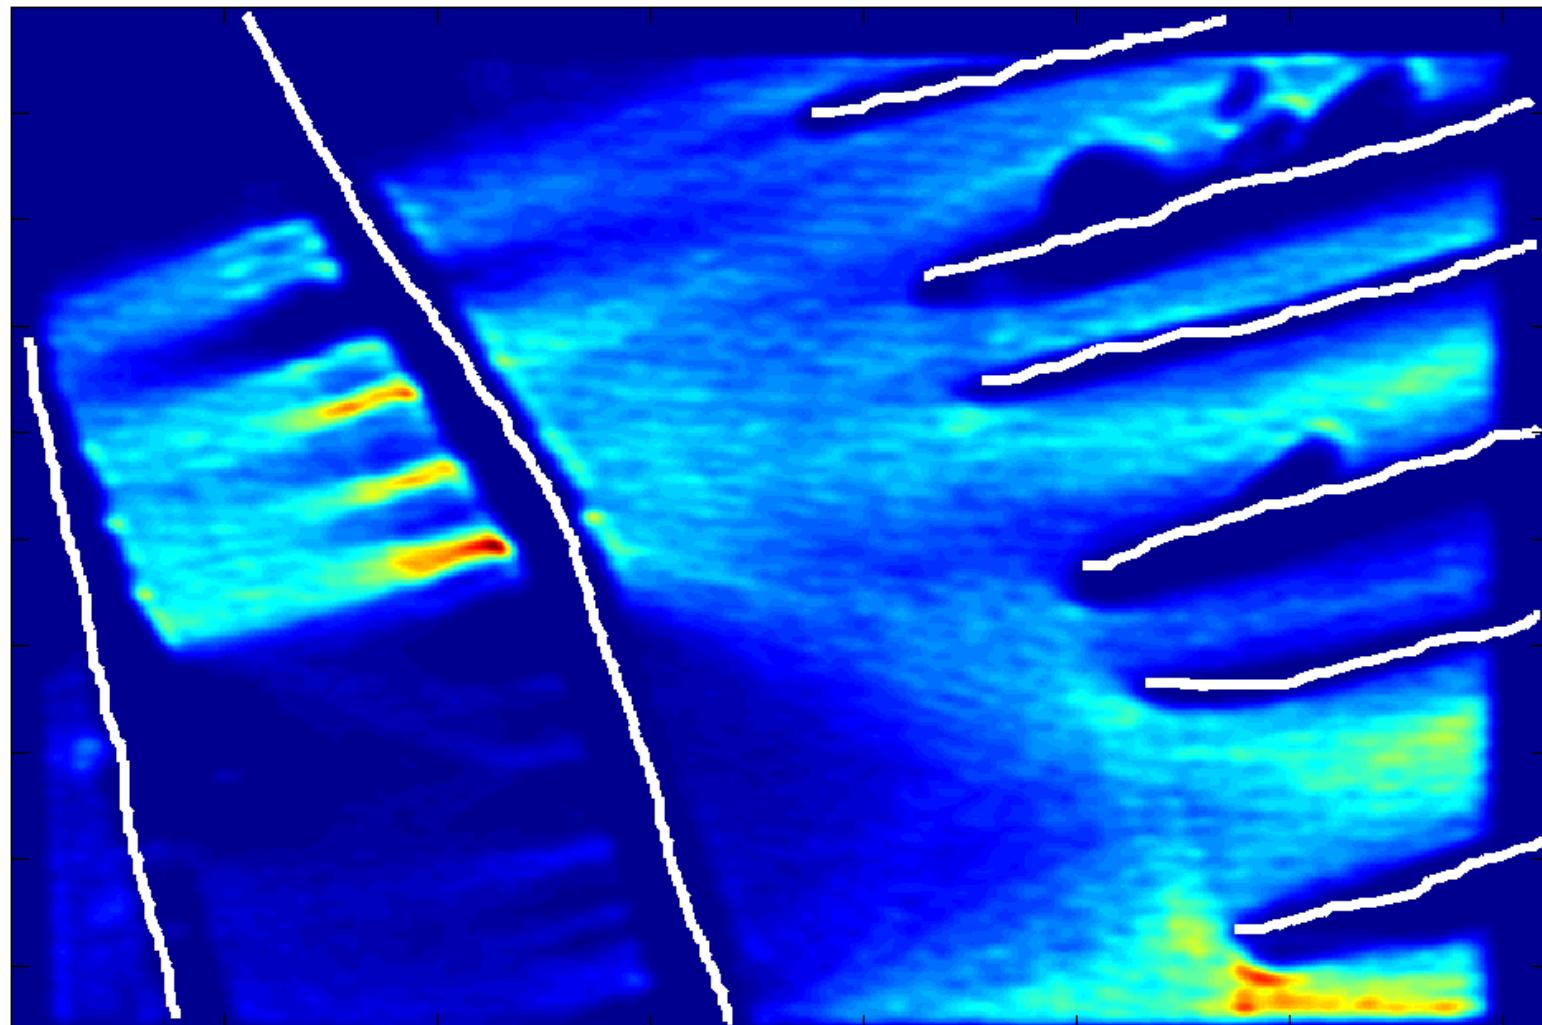
GateA Path Counts

movie



Maintain a running count of number of people whose
trajectories cross a set of user-specified lines (color-coded).

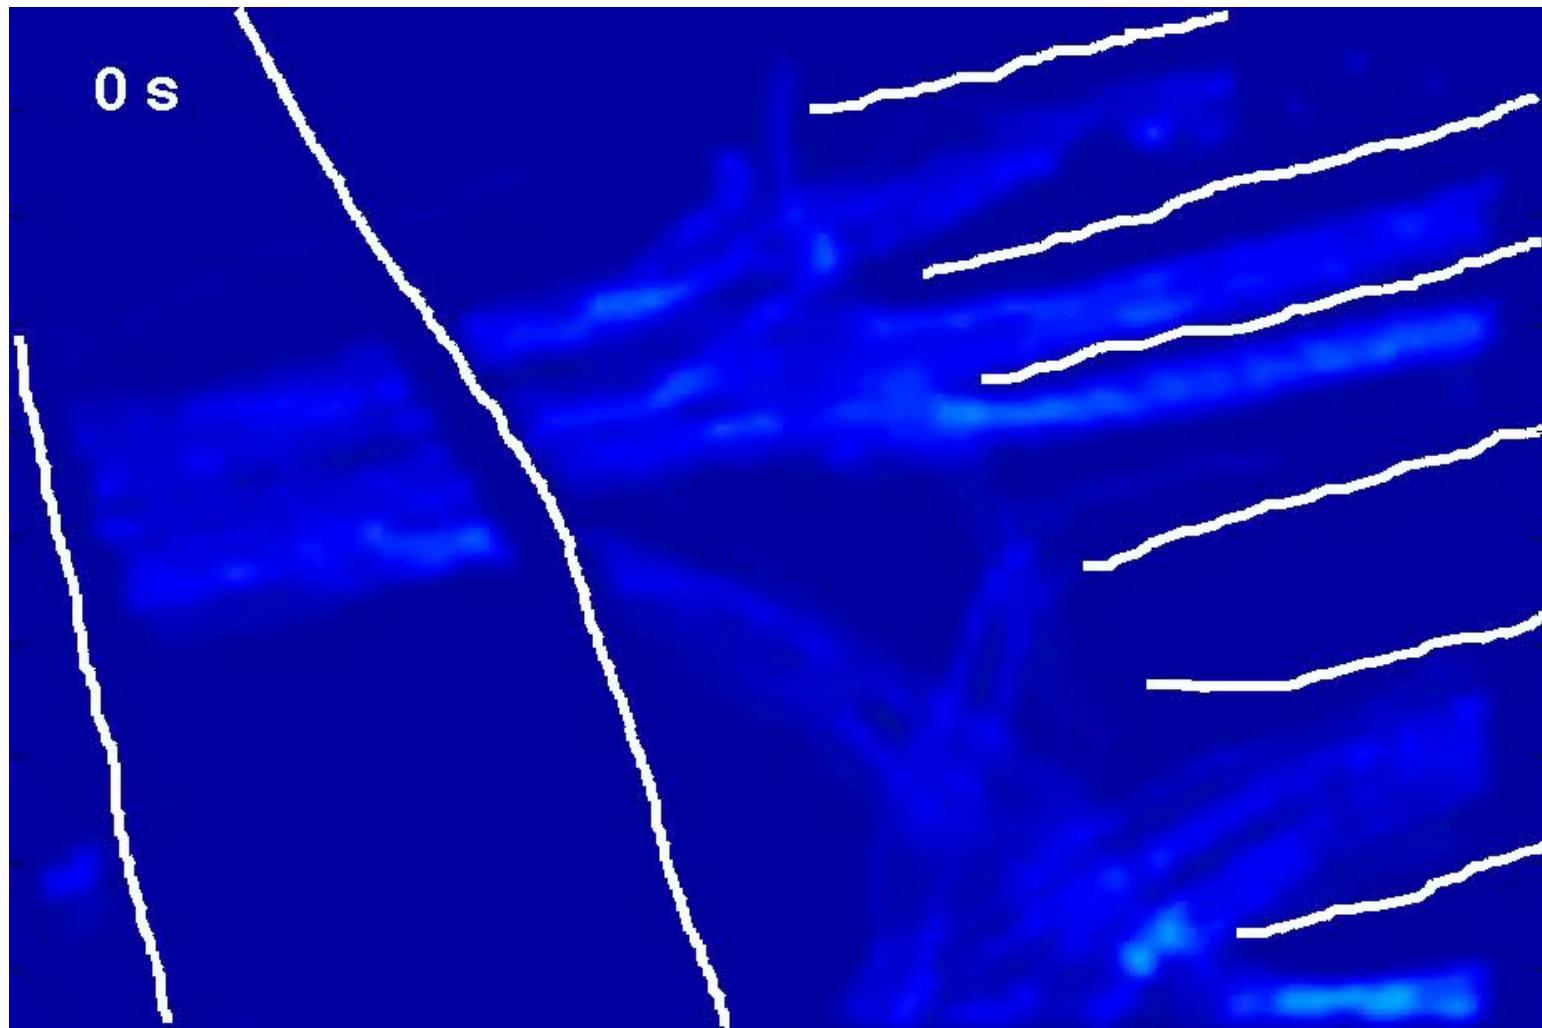
Crowd Flow/Density



30 minute period

Crowd Flow/Density

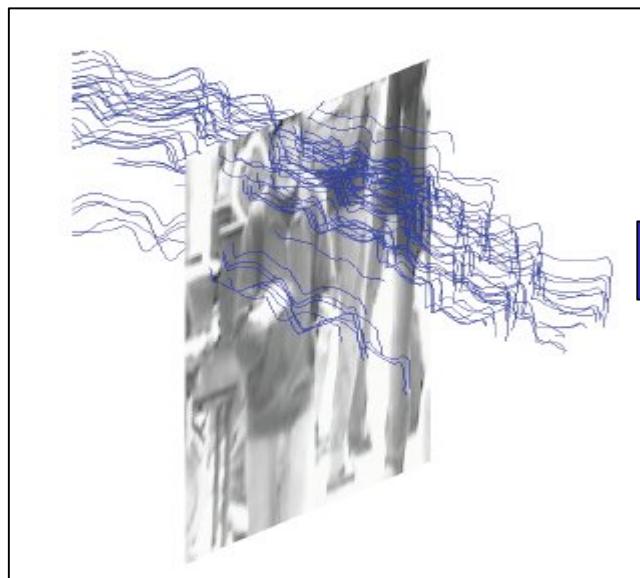
movie



Time Lapse. Integrated over spatial/temporal windows.

Motion Segmentation

Idea: track many small features (e.g. corners) over time and cluster sets of features that have similar motion.



corner trajectories

clustering

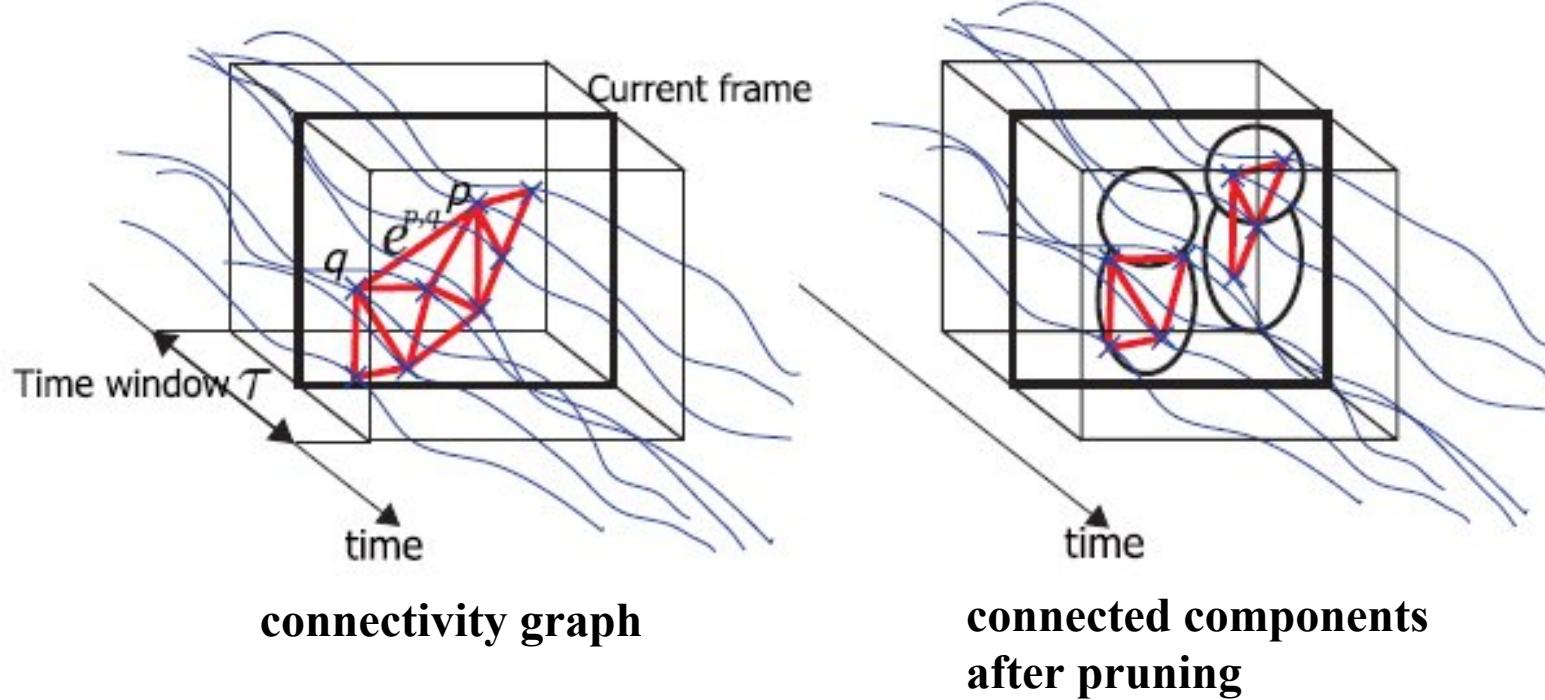
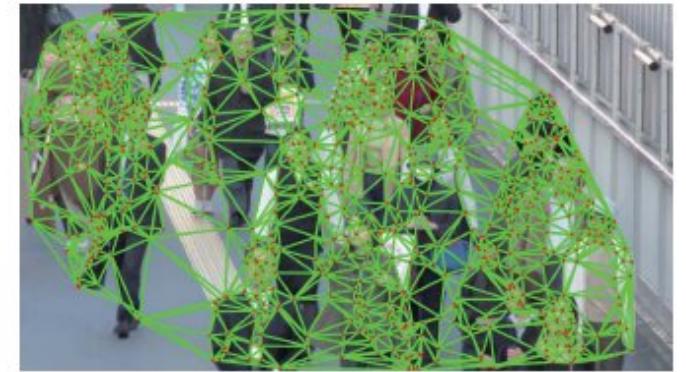


independently moving objects

- G. J. Brostow and R. Cipolla, “Unsupervised bayesian detection of independent motion in crowds,” in IEEE Conference on Computer Vision and Pattern Recognition, 2006, pp. 594–601.
- V. Rabaud and S. Belongie, “Counting crowded moving objects,” in IEEE Computer Vision and Pattern Recognition, New York City, 2006, pp. 705–711.
- D. Sugimura, K. Kitani, T. Okabe, Y. Sato, and A. Sugimoto, “Using individuality to track individuals: Clustering individual trajectories in crowds using local appearance and frequency trait,” in International Conference on Computer Vision, 2009, pp. 1467–1474.

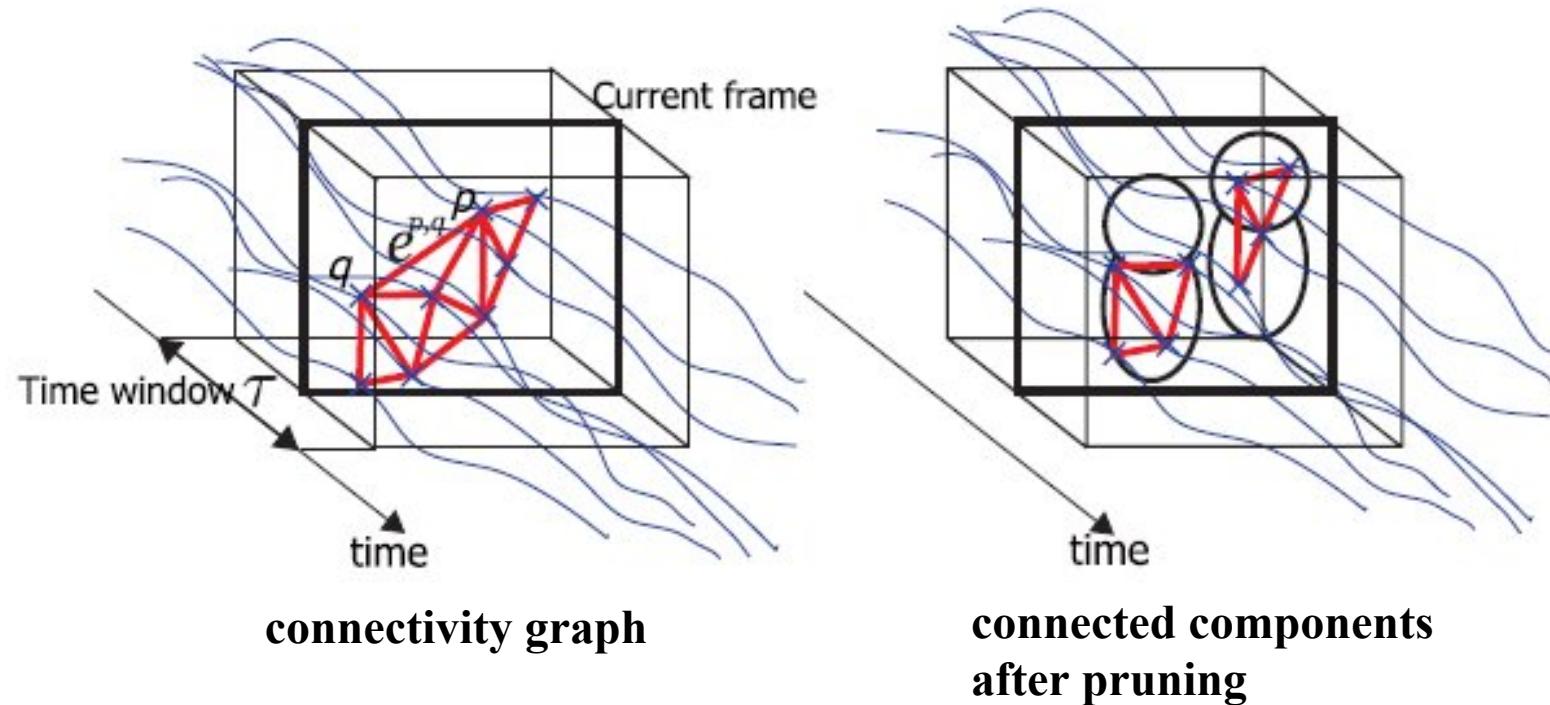
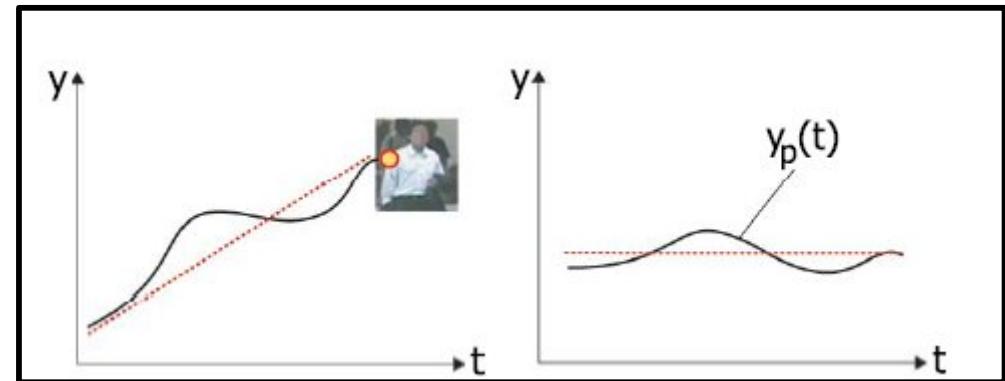
Motion Segmentation

Basic steps: Form a corner connectivity graph.
Assign each edge a dissimilarity score based on
distance and motion coherence of trajectories.
Prune edges with high scores. The remaining
connected components are the independent
objects.



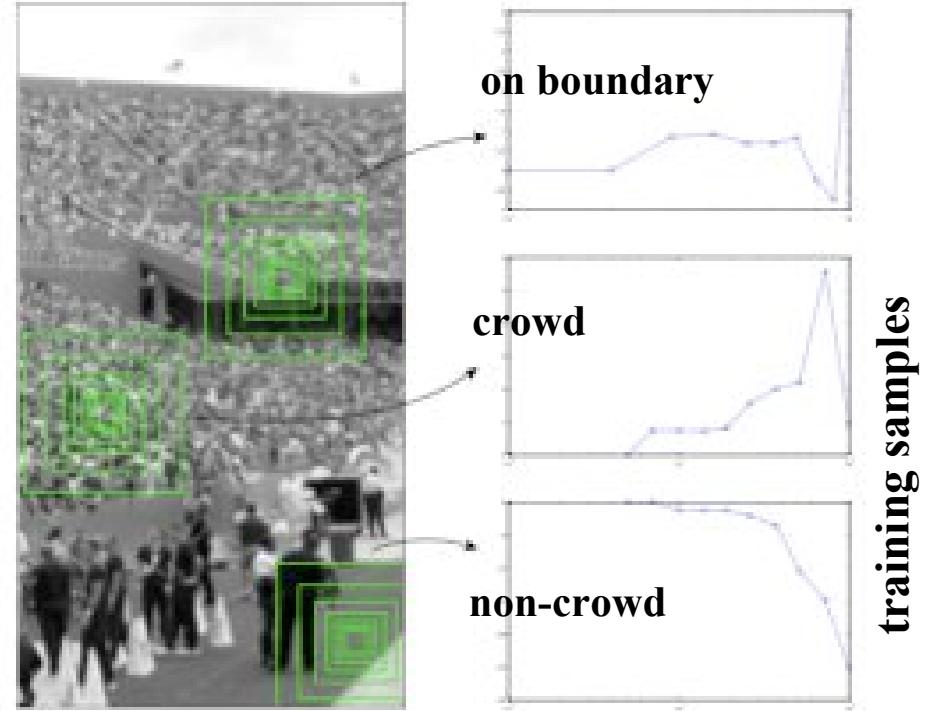
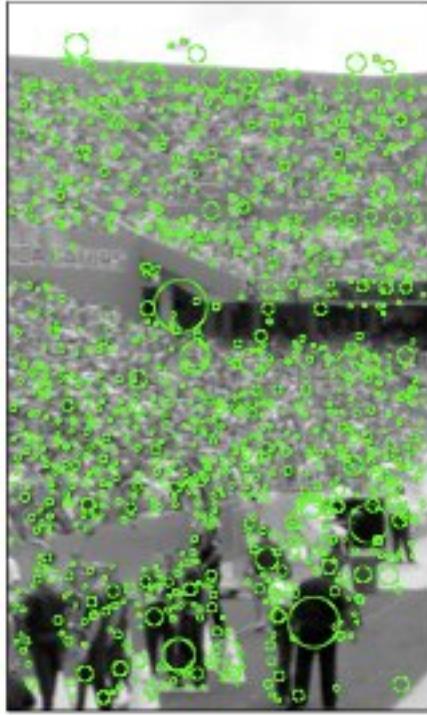
Motion Segmentation

Note: Sugimara et.al. add a feature based on gait periodicity to help disambiguate nearby people.



Texture-based Crowd Detection

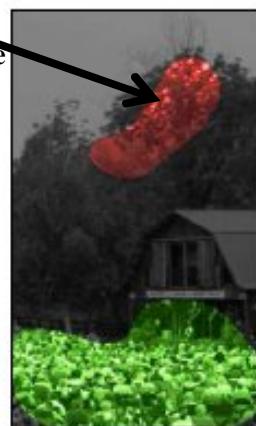
Arandjelovic, "Crowd Detection from Still Images," BMVC 2008



- SIFT descriptors
- K-means clustering to form “SIFT-Words”
- Likelihood ratio of distributions of word counts over 10 patch sizes yields 10-D feature vector
- Radial basis SVM for classification into crowd / non-crowd

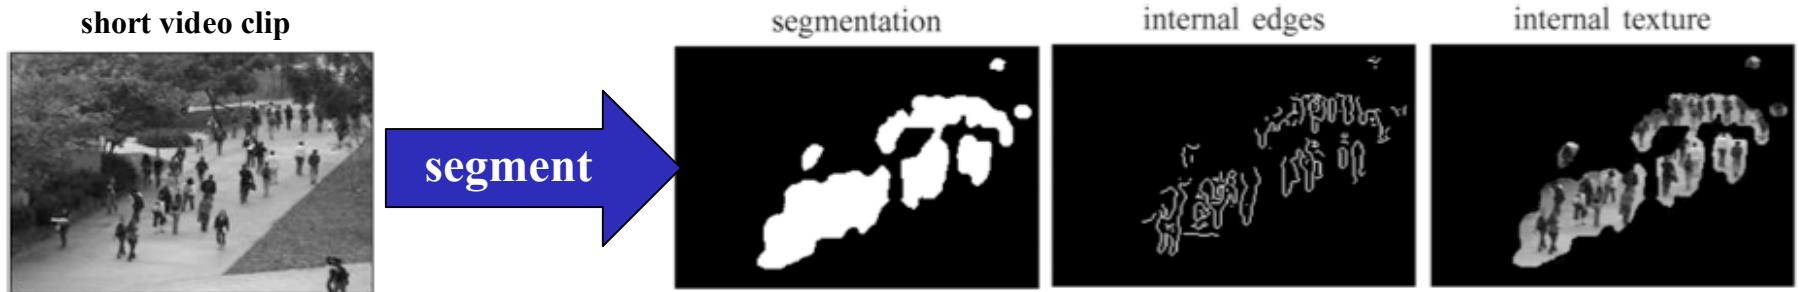
Texture-based Crowd Detection

Sparse classifications turned into dense segmentation using graph cuts. Unary costs based on SVM output and pairwise costs based on magnitude of patch likelihood scores (small magnitudes indicate interclass boundaries).



Texture-based Counting

Chan and Vasconcelos, “Counting People with Low-level Features and Bayesian Regression”,
IEEE Transactions on Image Processing, Vol 21 (4), 2160-2177, April 2012



**motion segmentation
using dynamic textures**

Extract feature vector for each frame:

- **region features**

e.g. area, perimeter, num connected components...

- **internal edge features**

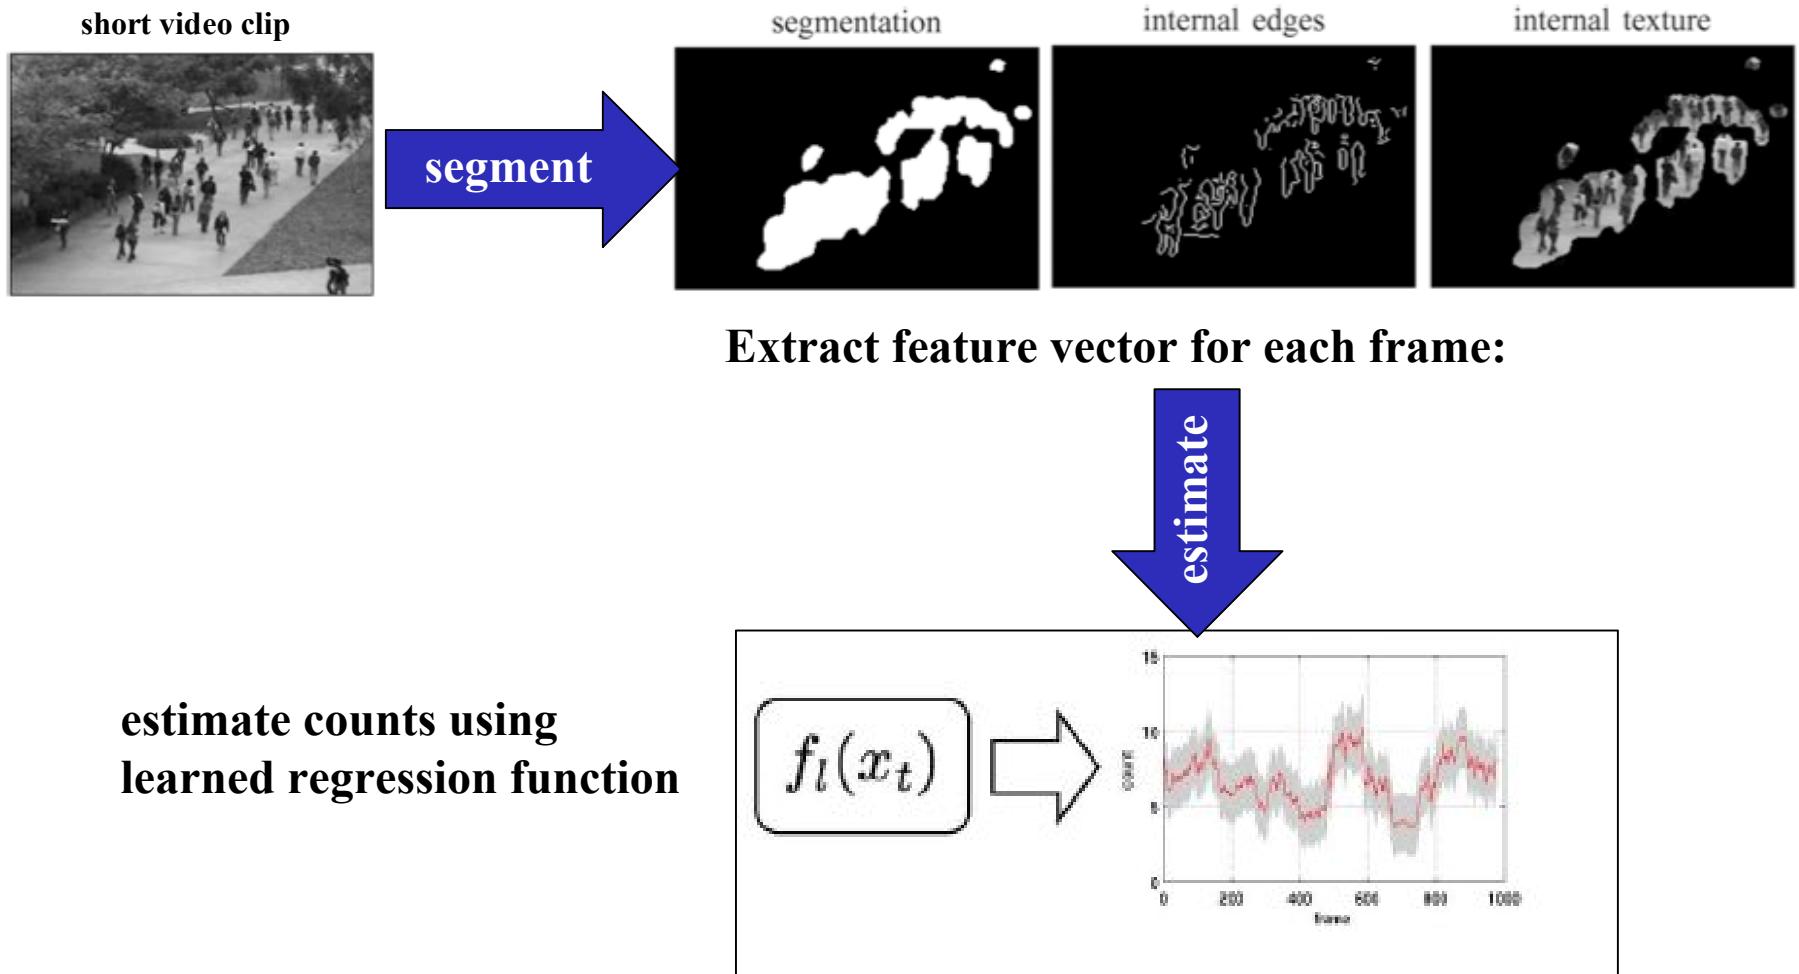
e.g. num edges, histogram of orientations

- **grey-level texture features**

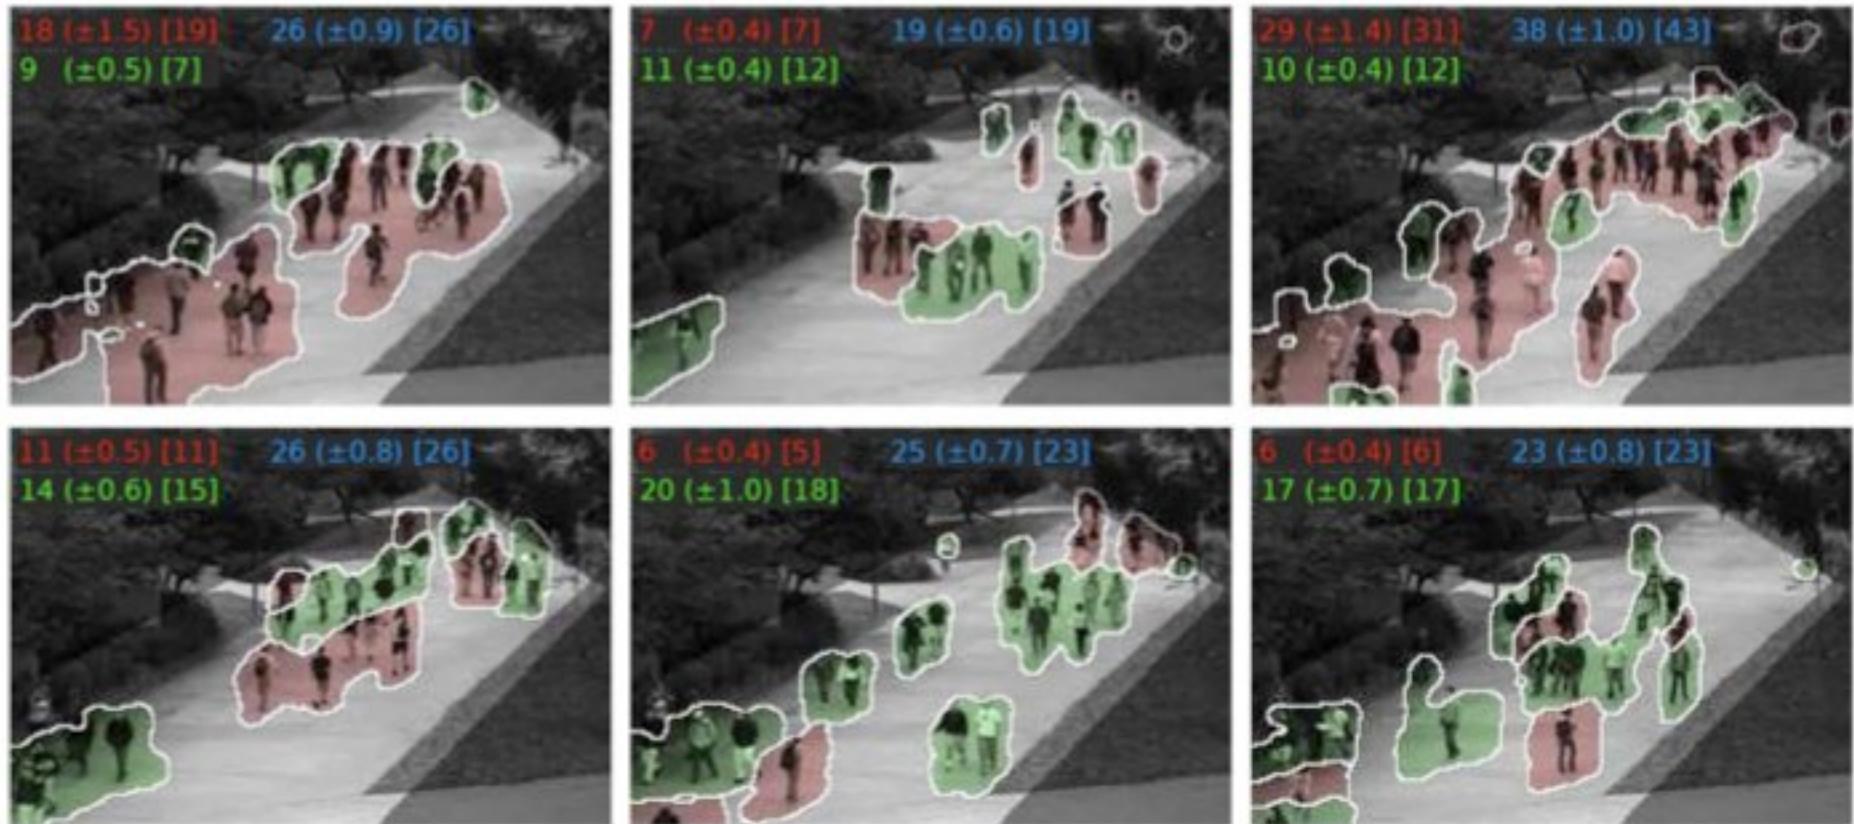
e.g. homogeneity, energy, entropy

Texture-based Counting

Chan and Vasconcelos, “Counting People with Low-level Features and Bayesian Regression”,
IEEE Transactions on Image Processing, Vol 21 (4), 2160-2177, April 2012

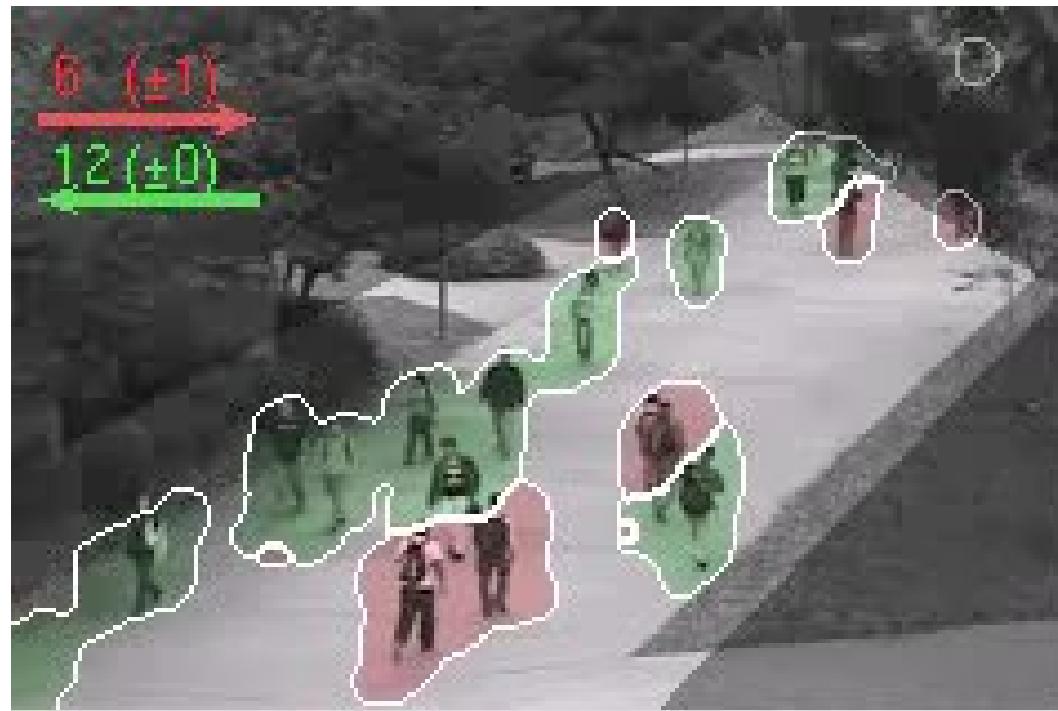


Texture-based Crowd Detection



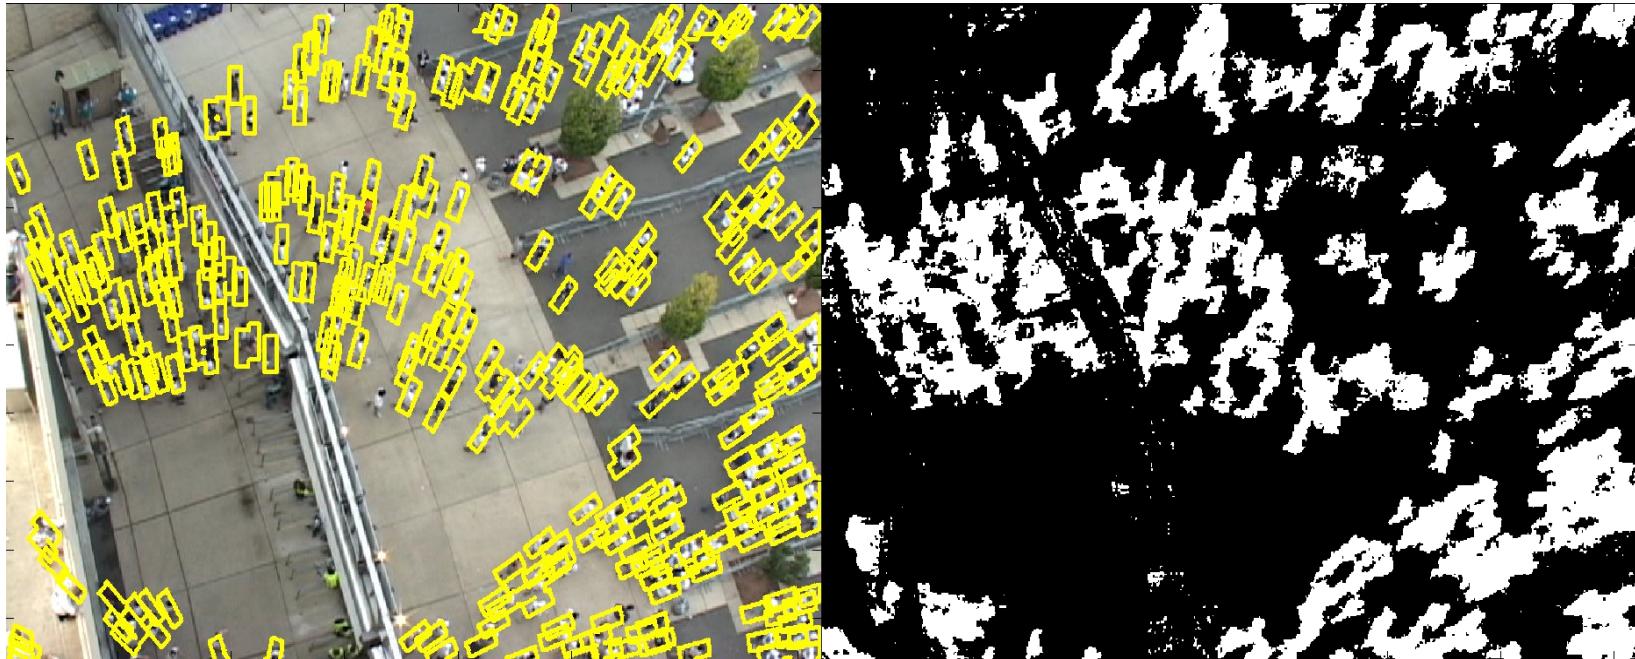
green/red = crowd walking towards/away blue = total
numeric results formatted as: estimated count (uncertainty) [true count]

Texture-based Crowd Detection



green/red = crowd walking towards/away
numeric results formatted as: estimated count (uncertainty)

Our Approach*

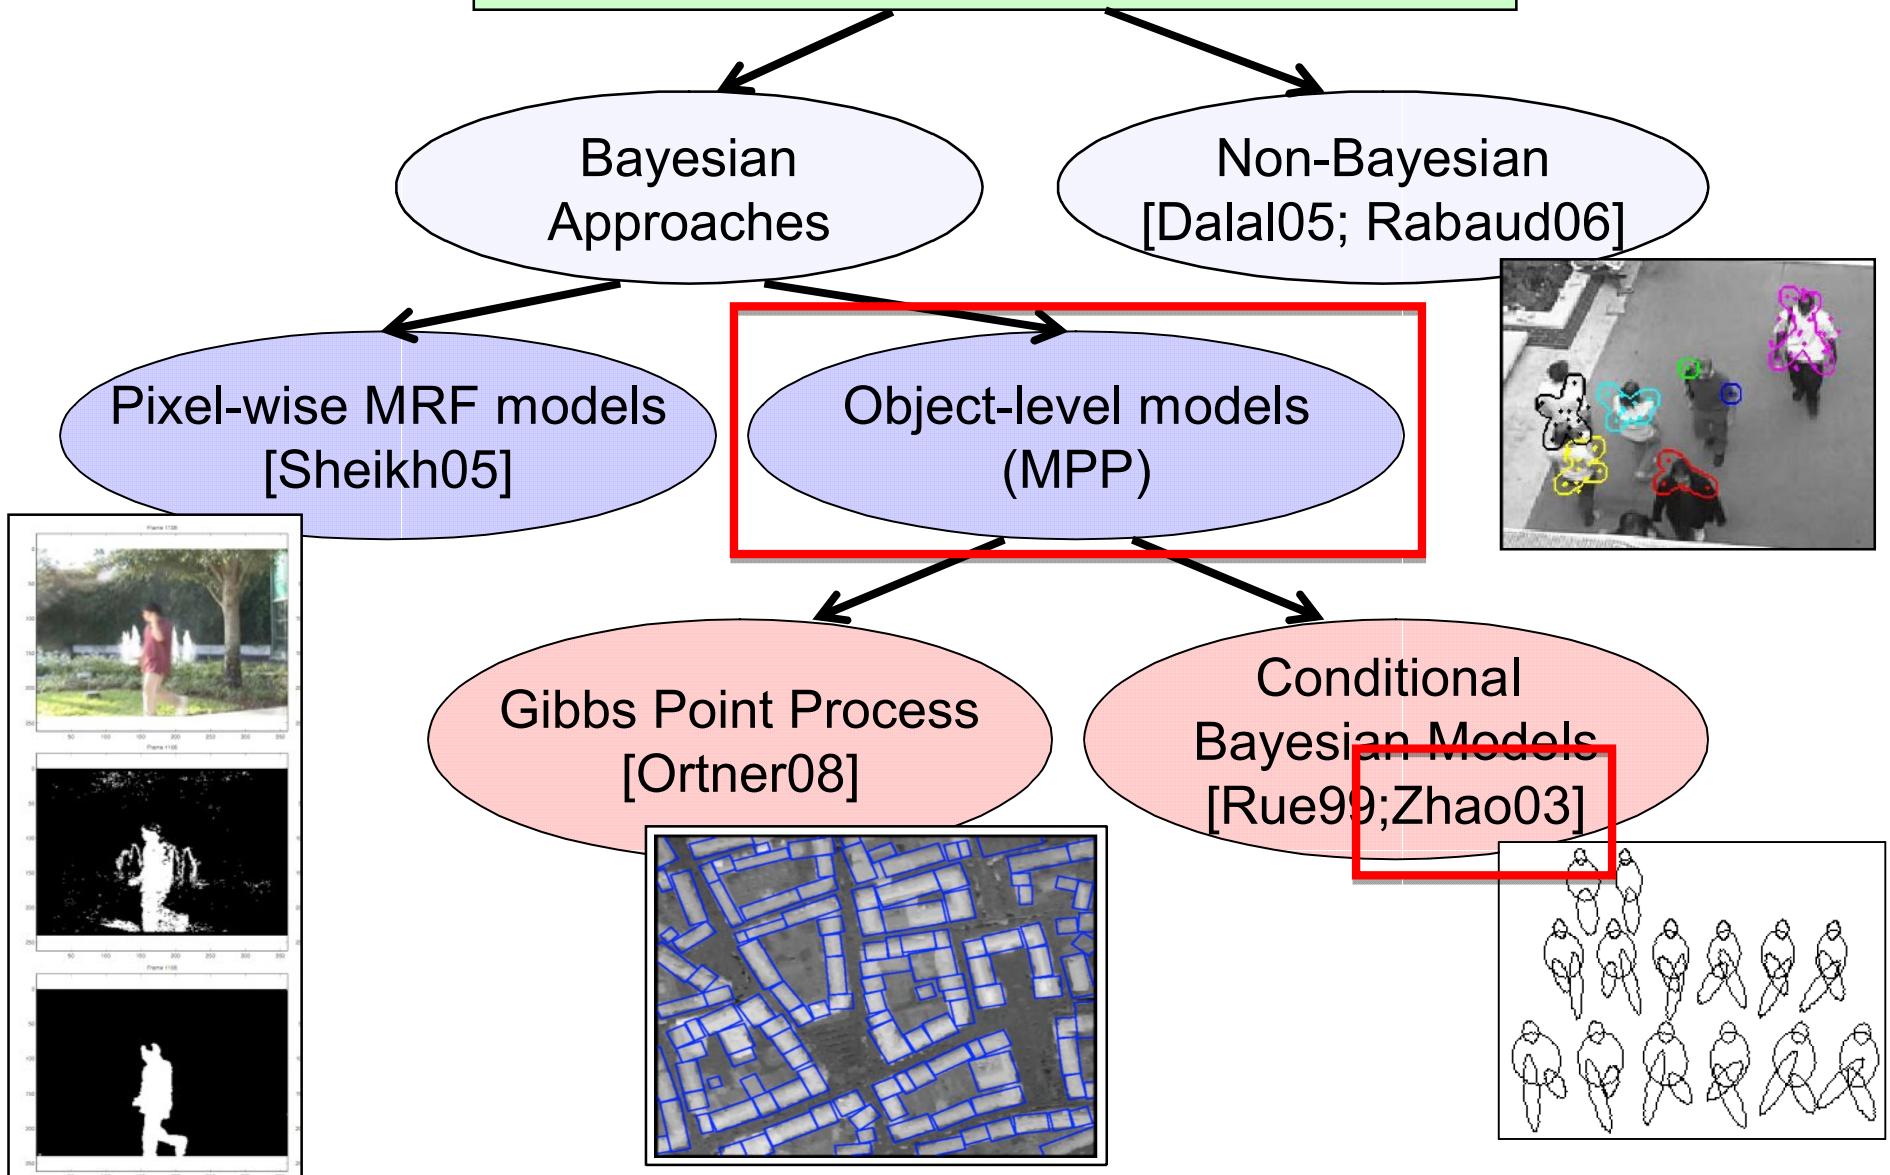


Bayesian Marked Point Process

- the prior models expected size/shape of people, indexed by location
- the likelihood measures how well a proposed configuration explains the data
- MAP estimate of number/configuration of people is found using RJ-MCMC

*Weina Ge and R.Collins, “Marked Point Processes for Crowd Counting,”
IEEE Computer Vision and Pattern Recognition, Miami, FL, June 2009.

Foreground Object Detection

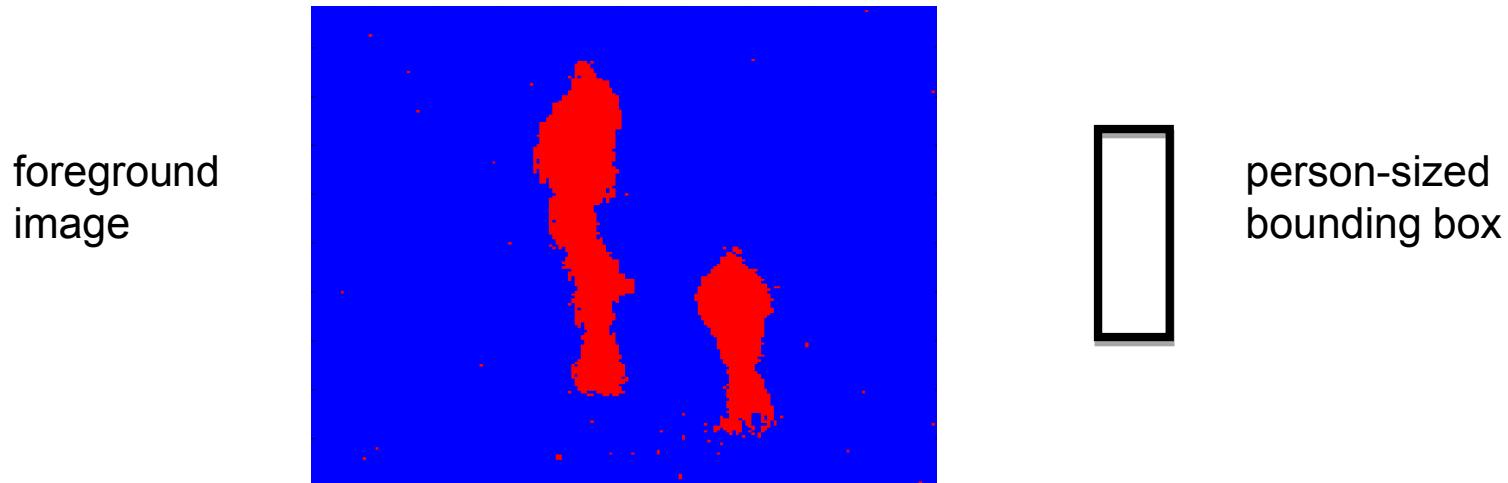


Marked Point Process

- Spatial Point Process
 - Distribution of a set of points in a bounded space
- Marked Point Process (MPP)
 - A spatial point process + a “mark” process
- Examples
 - Spatial process could model tree locations in a forest and the mark could be tree height
 - Spatial process could model cell locations on a microscope slide and the mark could be a polygonal representation of the cell boundary [Rue and Hurn, 1999]

Simplified Problem Statement

Given a foreground image, find a configuration of bounding boxes* that cover a majority of foreground pixels while leaving a majority of background pixels uncovered.



As an MPP:

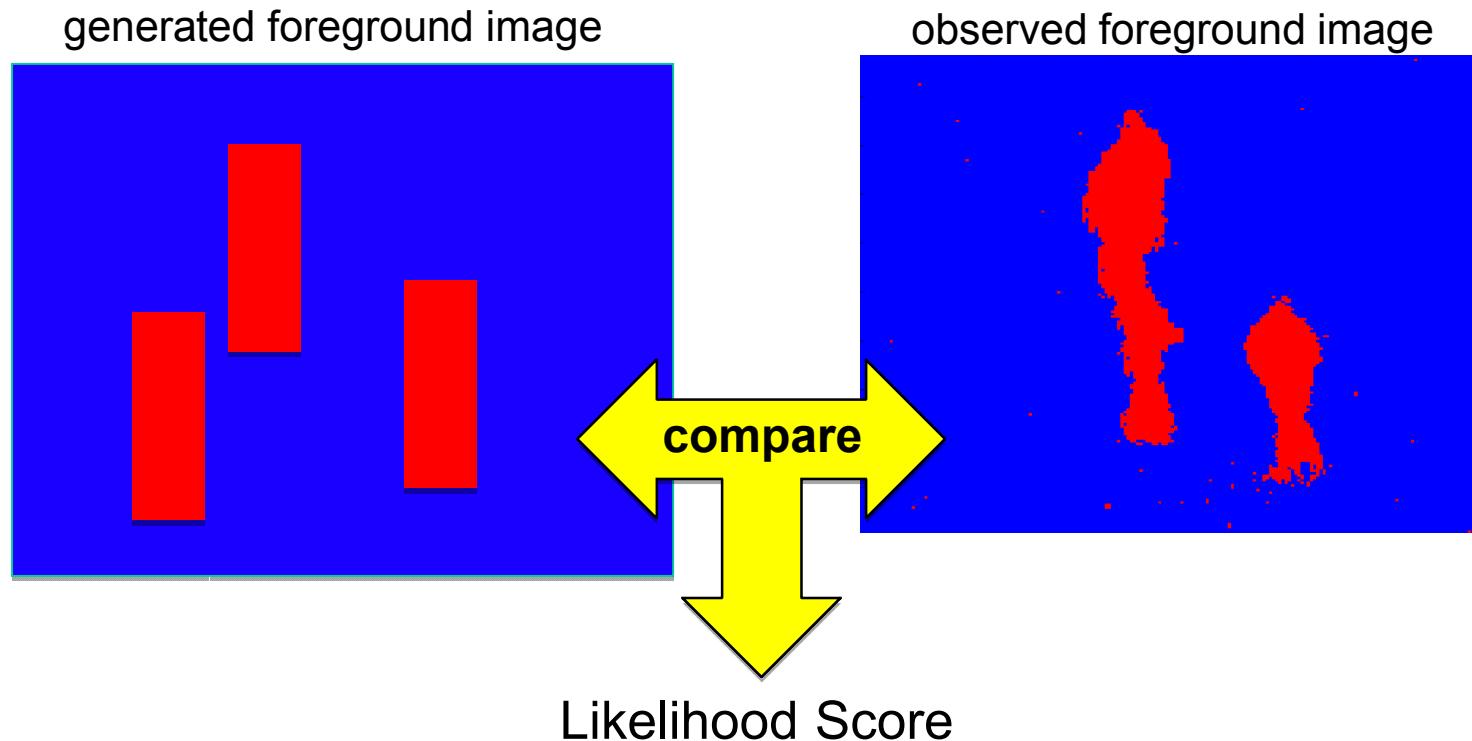
- Spatial process models number and (x,y) locations of bounding boxes
- Mark process models (height,width,orientation) of each bounding box.

*Footnote: We will add more realistic “shape” models in a moment

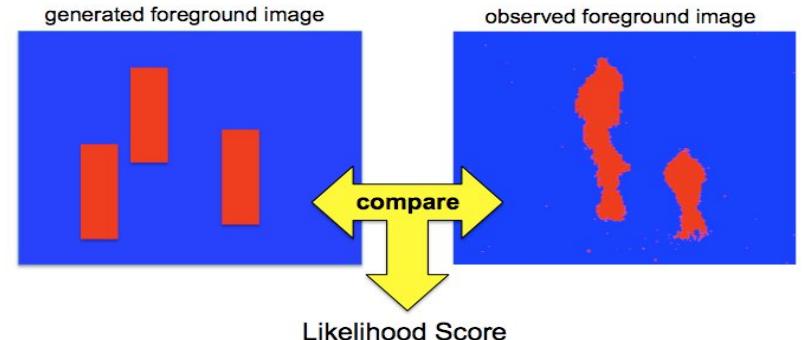
Likelihood Score

To measure how “good” a proposed configuration is, we generate a foreground image from it and compare with the observed foreground image to get a likelihood score.

config = $\{\{x_1, y_1, w_1, h_1, \text{theta}_1\}, \{x_2, y_2, w_2, h_2, \text{theta}_2\}, \{x_3, y_3, w_3, h_3, \text{theta}_3\}\}$



Likelihood Score



Bernoulli distribution

model

$p_{00} = p(y_i = 0|x_i = 0)$ = prob of observing background given a label of background

$p_{01} = p(y_i = 0|x_i = 1)$ = prob of observing background given a label of foreground

$p_{10} = p(y_i = 1|x_i = 0)$ = prob of observing foreground given a label of background

$p_{11} = p(y_i = 1|x_i = 1)$ = prob of observing foreground given a label of foreground

c_{00} = count of pixels where observation is background and label is background

c_{01} = count of pixels where observation is background and label is foreground

c_{10} = count of pixels where observation is foreground and label is background

c_{11} = count of pixels where observation is foreground and label is foreground

likelihood

$$L(Y|X) = \prod_{i=1}^N p(y_i|x_i) = p_{00}^{c_{00}} p_{01}^{c_{01}} p_{10}^{c_{10}} p_{11}^{c_{11}}$$

simplify, by assuming

$$p_{00} = p_{11} = \mu \quad \text{and} \quad p_{01} = p_{10} = 1 - \mu$$

Number of pixels
that disagree!

log likelihood

$$\begin{aligned} \log L(Y|X) &= (c_{00} + c_{11}) \log \mu + (c_{01} + c_{10}) \log(1 - \mu) \\ &= [N - (c_{01} + c_{10})] \log \mu + (c_{01} + c_{10}) \log(1 - \mu) \\ &= N \log \mu - (c_{01} + c_{10}) [\log \mu - \log(1 - \mu)] \end{aligned}$$

Prior Model

We use a prior to model our expectations about bounding box configurations in the image

$$\pi(\mathbf{o}_i) = \underbrace{\pi(p_i)}_{\text{Prior for bounding box } i} \underbrace{\pi(w_i, h_i, \theta_i | p_i)}_{\text{Prior on location (center point)}} \underbrace{\pi(w_i, h_i, \theta_i | p_i)}_{\text{Prior on bounding box height, width and orientation, conditioned on center location.}}$$

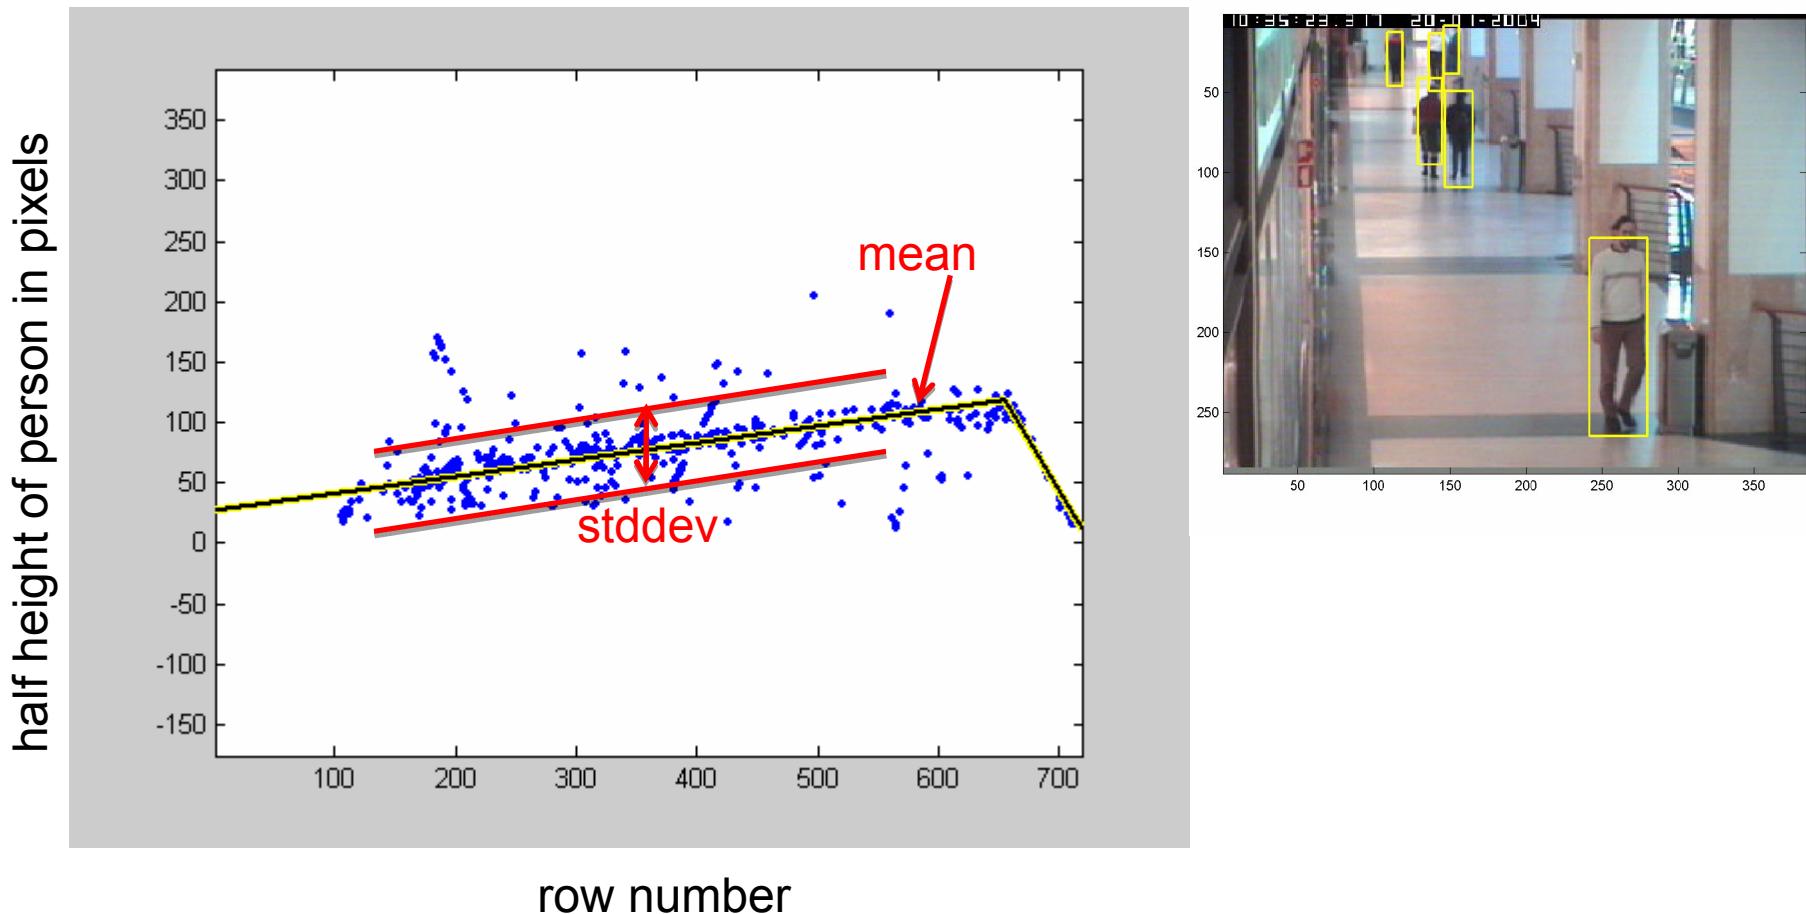
Prior for
bounding
box i

Prior on location
(center point)

Prior on bounding box
height, width and orientation,
conditioned on center location.

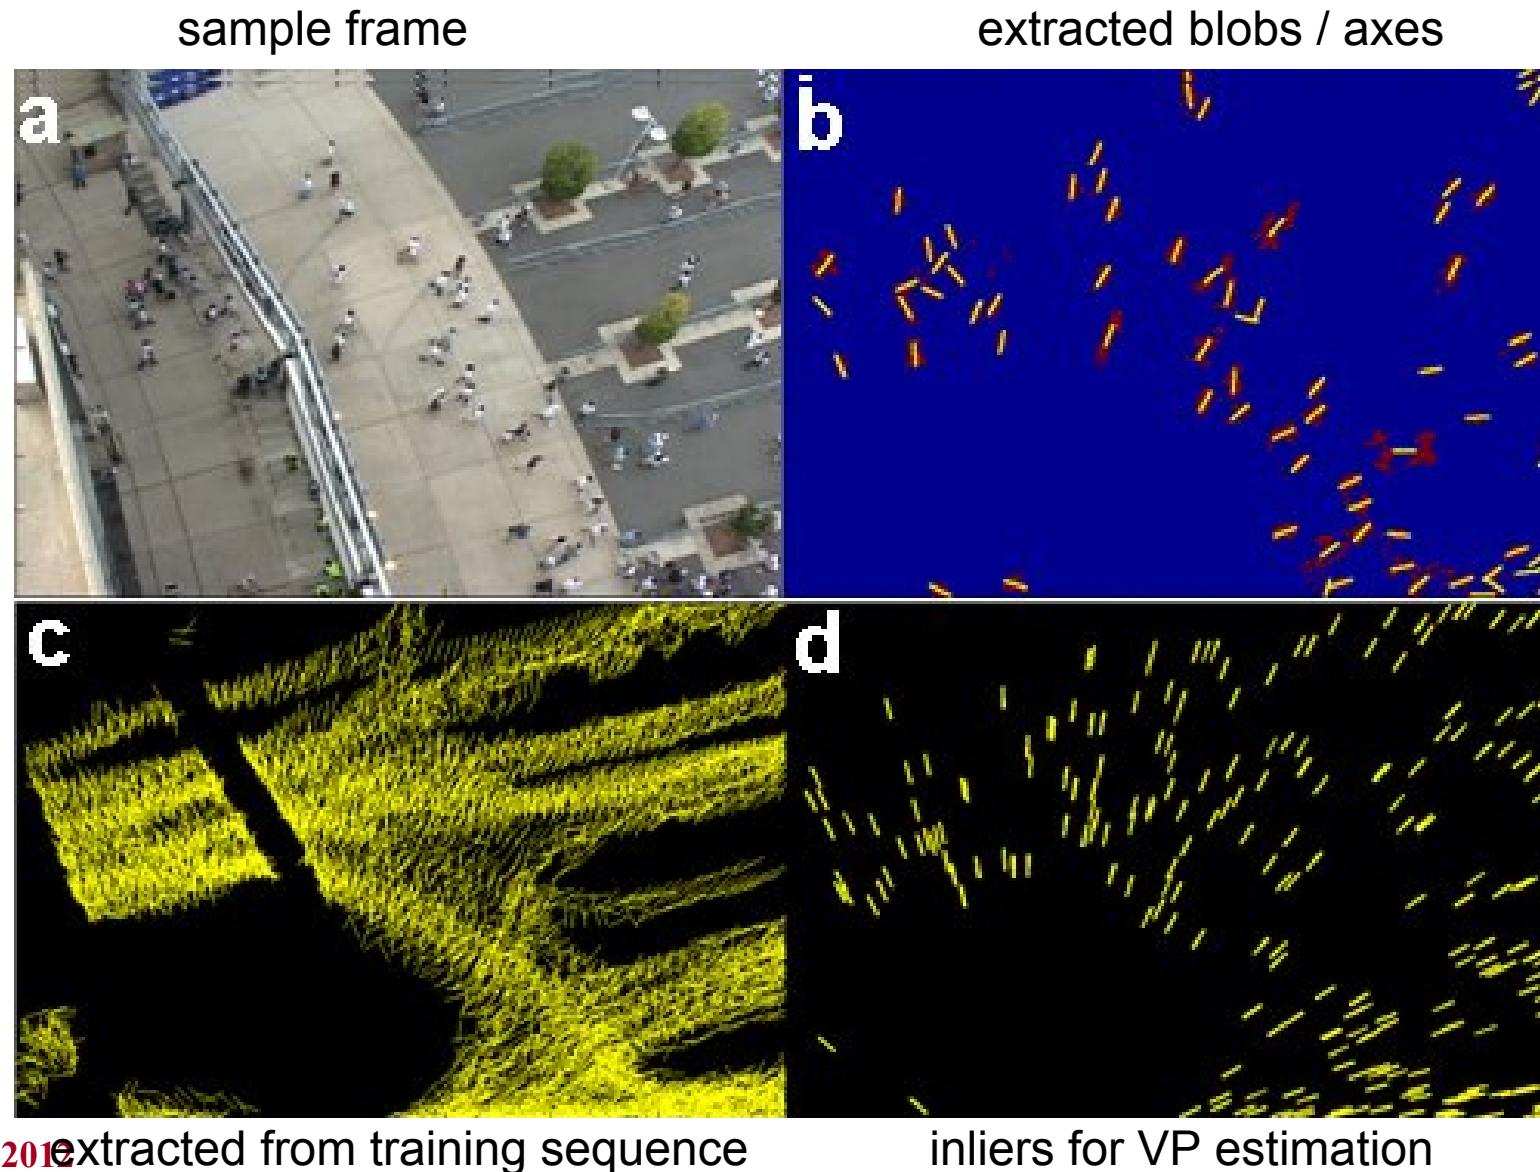
Estimating Priors

Example: learning height distribution as a function of image row



Estimating Priors

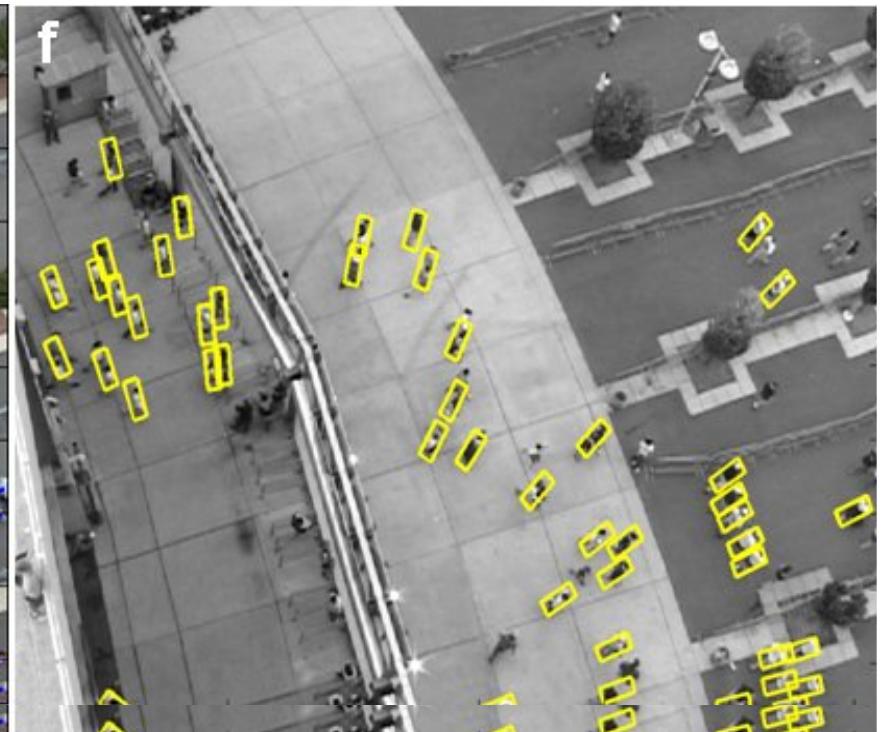
Example: learning orientation as a function of image location



Estimating Priors



estimated vanishing point



Scaled, oriented rectangles

Bottom line: it is not difficult to estimate priors on location, size and orientation of people as seen from a specific viewpoint.

Searching for the Max

The space of configurations is very large. We can't exhaustively search for the max likelihood configuration. We can't even really uniformly sample the space to a reasonable degree of accuracy.

$$\text{config}_k = \{\{x_1, y_1, w_1, h_1, \theta_1\}, \{x_2, y_2, w_2, h_2, \theta_2\}, \dots, \{x_k, y_k, w_k, h_k, \theta_k\}\}$$

Let N = number of possible locations for (x_i, y_i) in a k -person configuration.

Size of config_k = N^k

And we don't even know how many people there are...

Size of config space = $N^0 + N^1 + N^2 + N^3 + \dots$

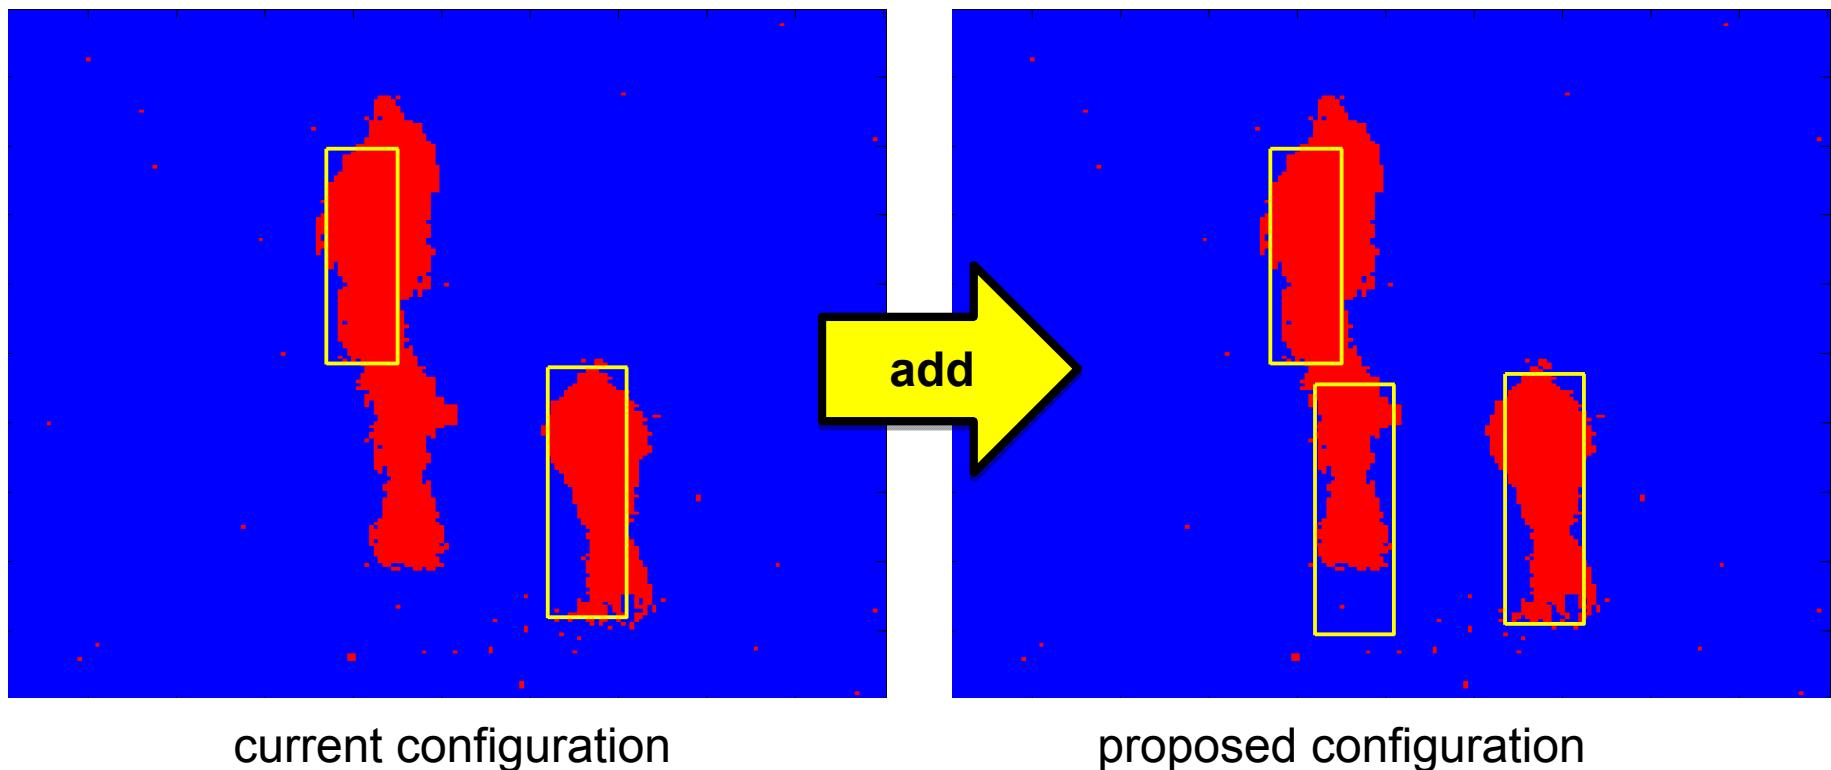
If we also wanted to search for width, height and orientation, this space would be even more huge.

Searching for the Max

- Local Search Approach
 - Given a current configuration, propose a small change to it
 - Compare likelihood of proposed config with likelihood of the current config
 - Decide whether to accept the change

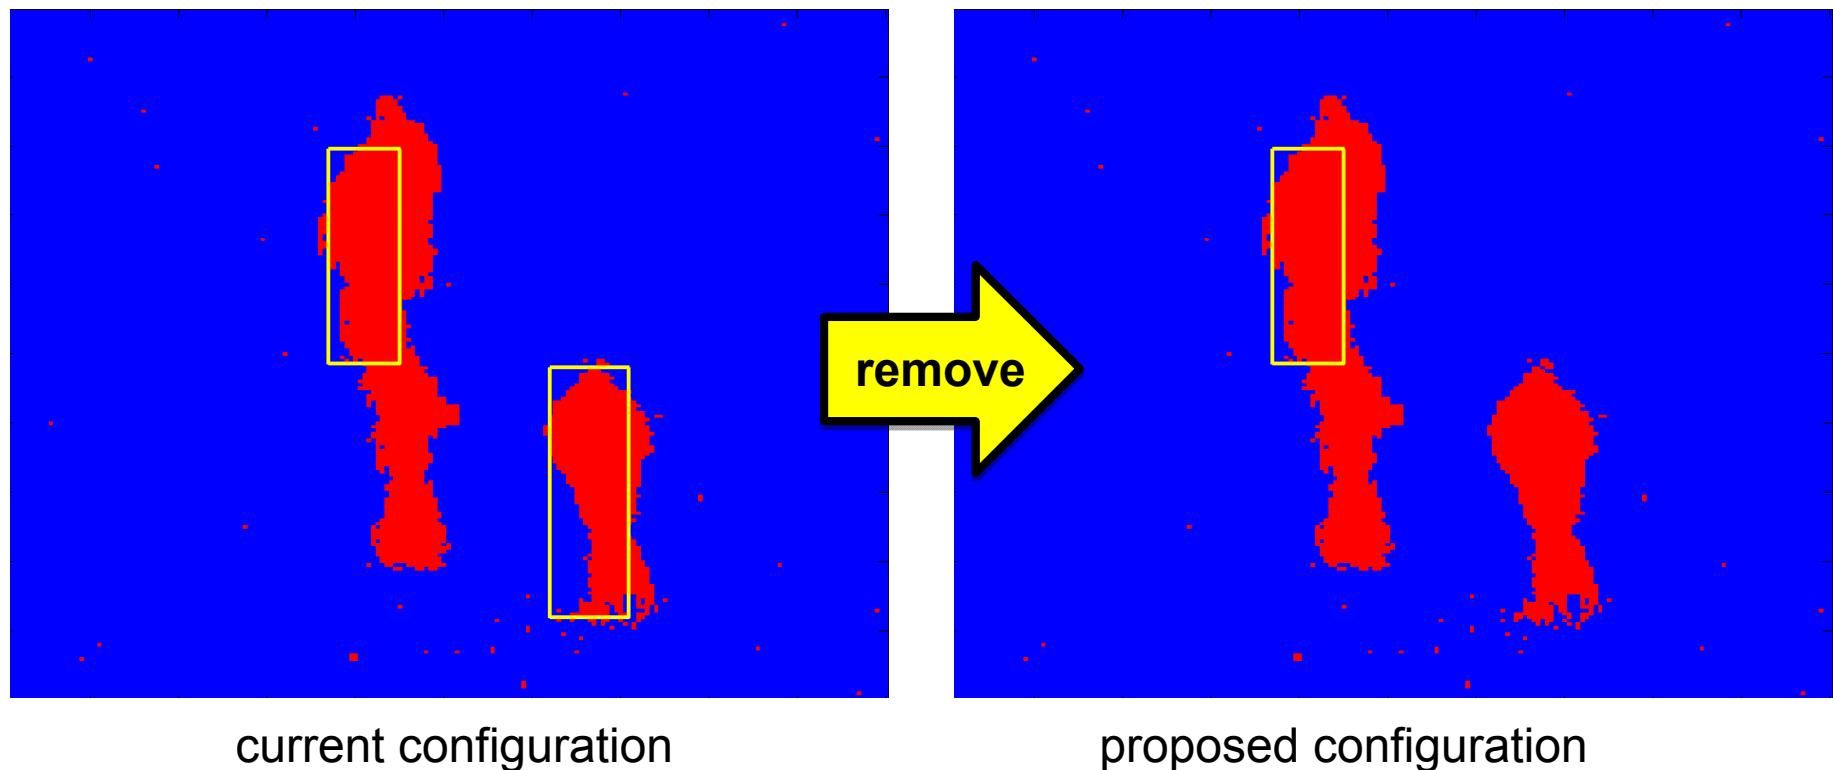
Proposals

- Add a rectangle (birth)



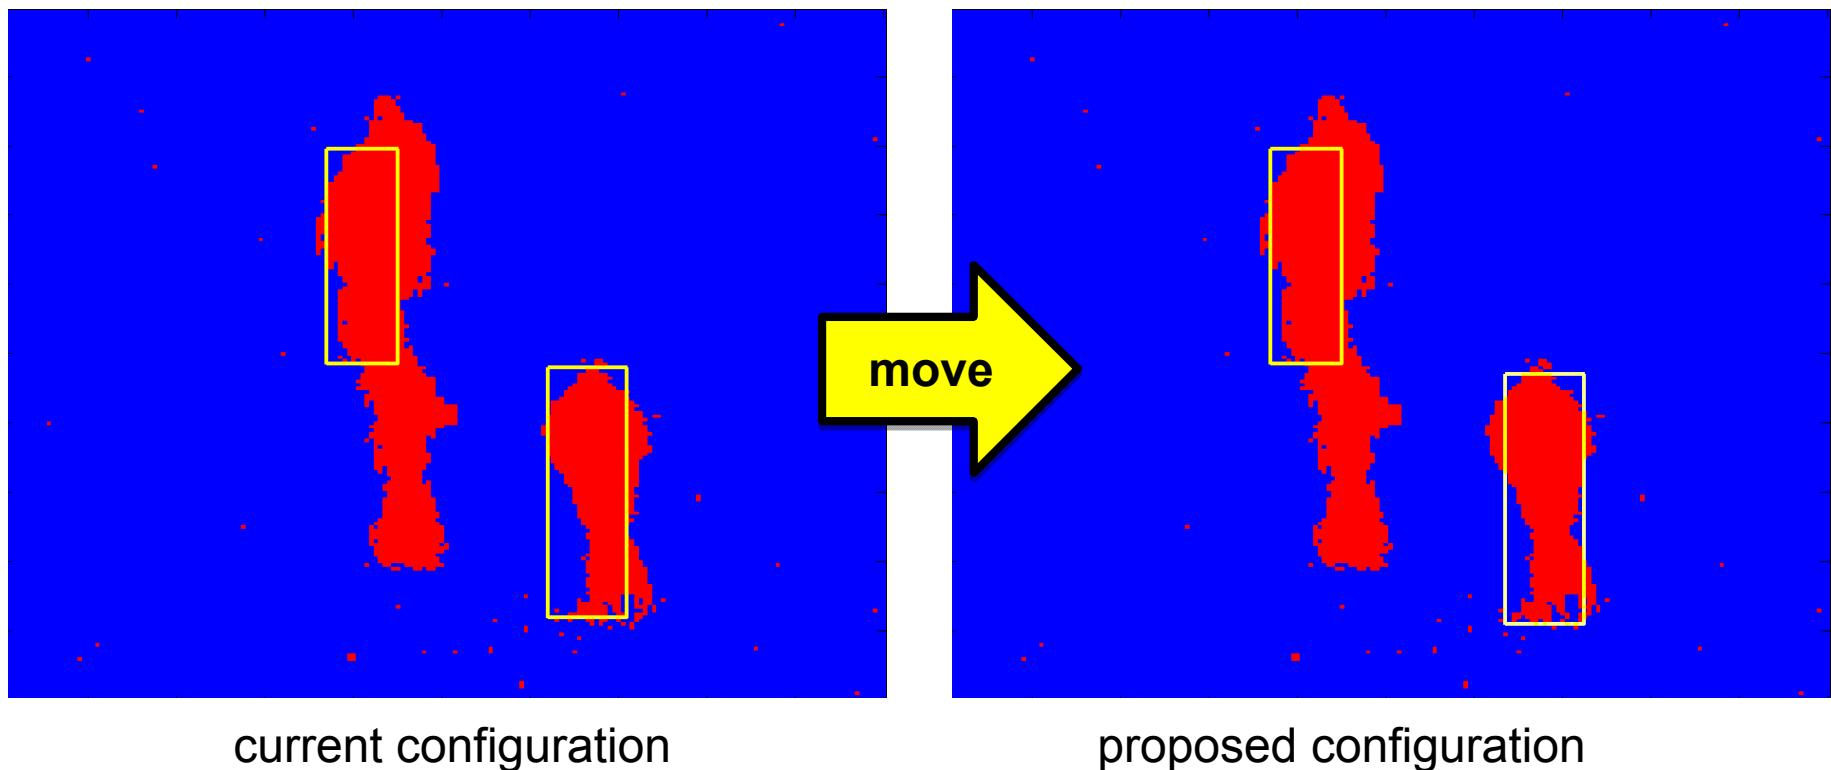
Proposals

- Remove a rectangle (death)



Proposals

- Move a rectangle



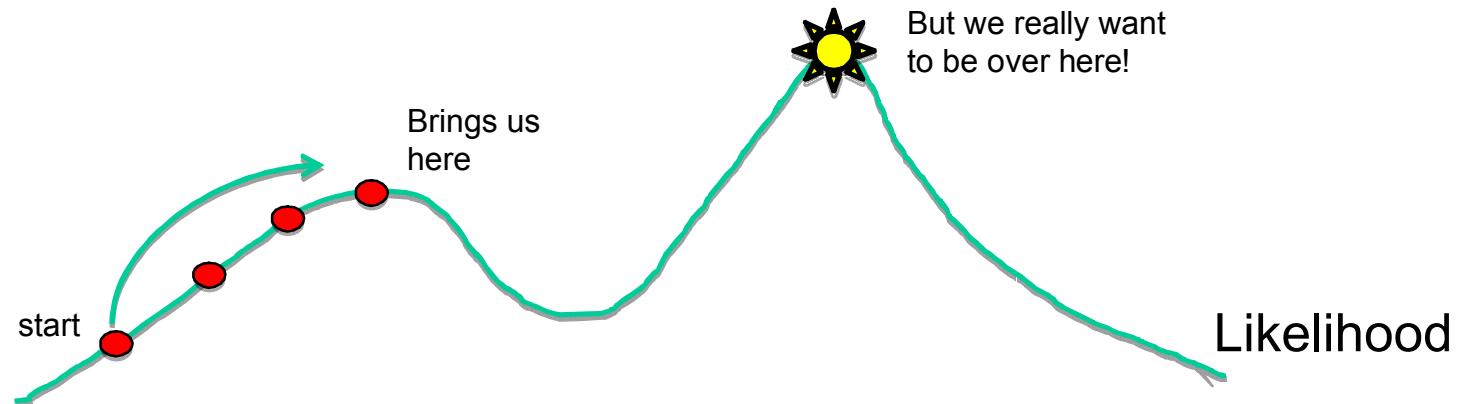
Searching for the Max

- Naïve Acceptance
 - Accept proposed configuration if it has a larger likelihood score, i.e.

$$\text{Compute } a = \frac{L(\text{proposed})}{L(\text{current})}$$

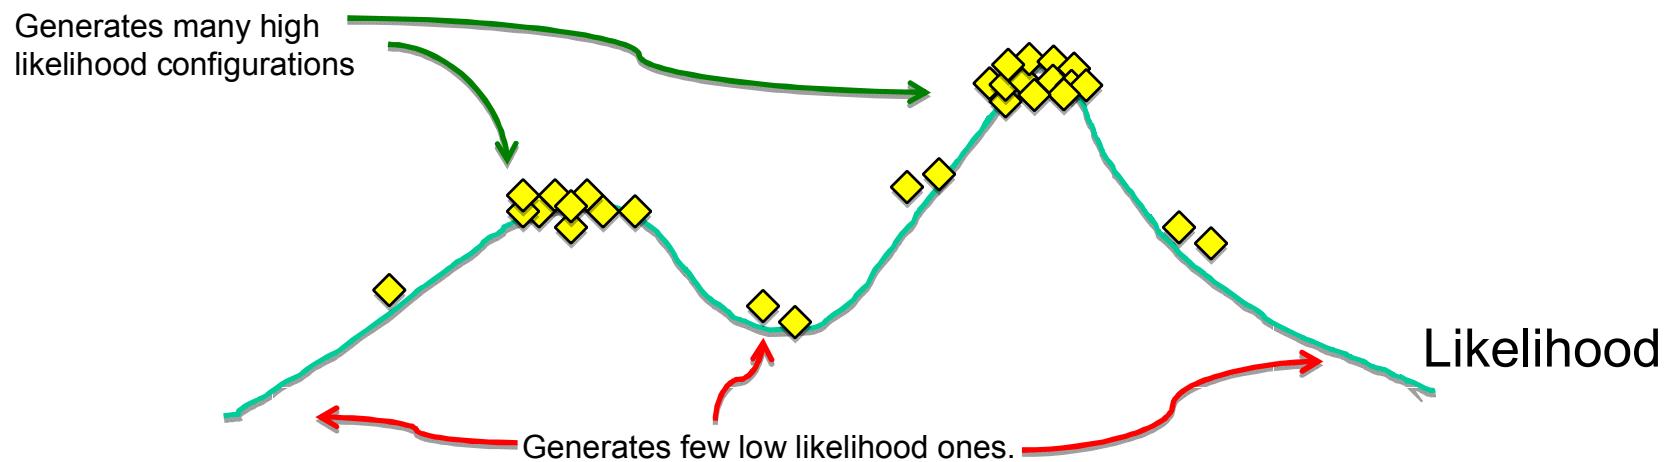
Accept if $a > 1$

- Problem: leads to hill-climbing behavior that gets stuck in local maxima



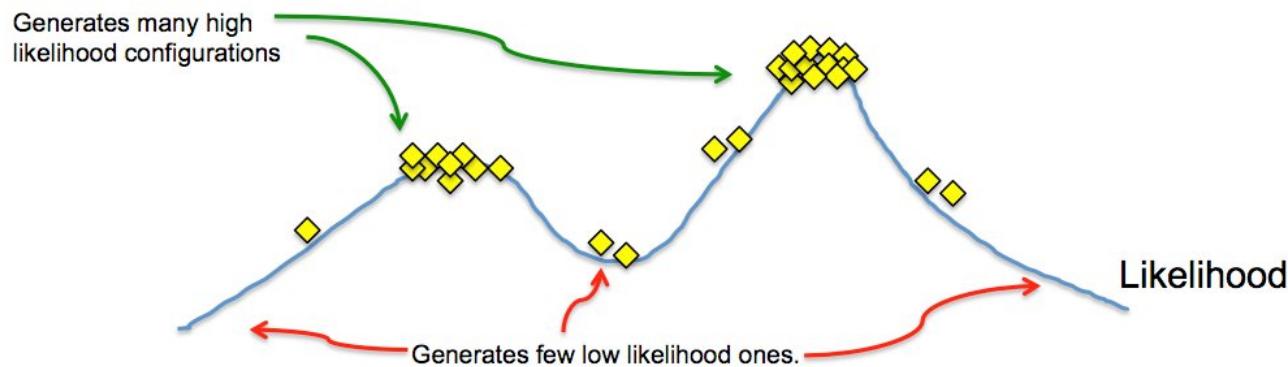
Searching for the Max

- The MCMC approach
 - Generate random configurations from a distribution proportional to the likelihood!



Searching for the Max

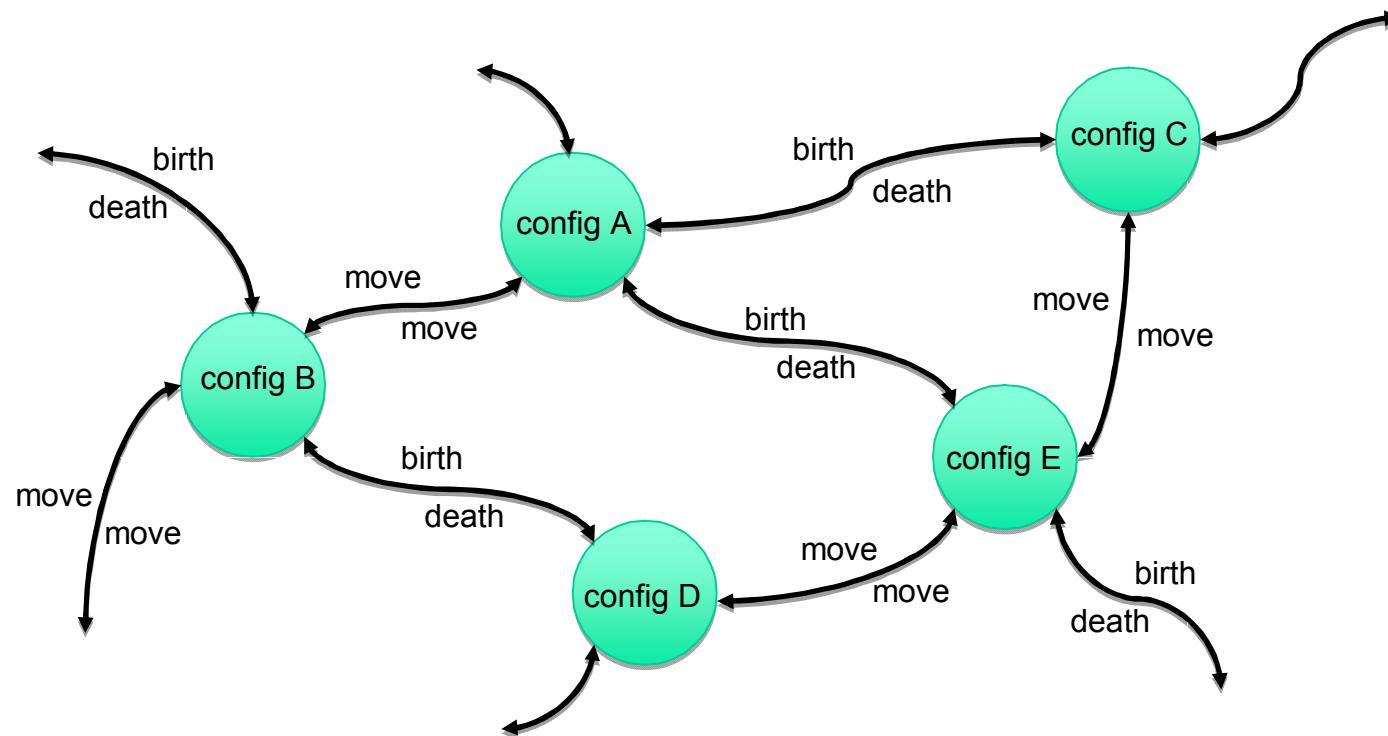
- The MCMC approach
 - Generate random configurations from a distribution proportional to the likelihood!



- This searches the space of configurations in an efficient way.
- Now just remember the generated configuration with the highest likelihood.

Sounds good, but how to do it?

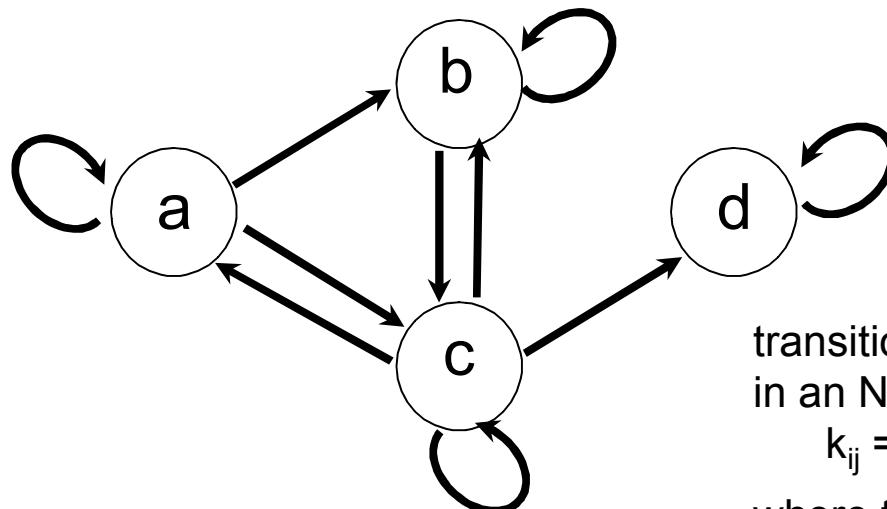
- Think of configurations as nodes in a graph.
- Put a link between nodes if you can get from one config to the other in one step (birth, death, move)



Recall: Markov Chains

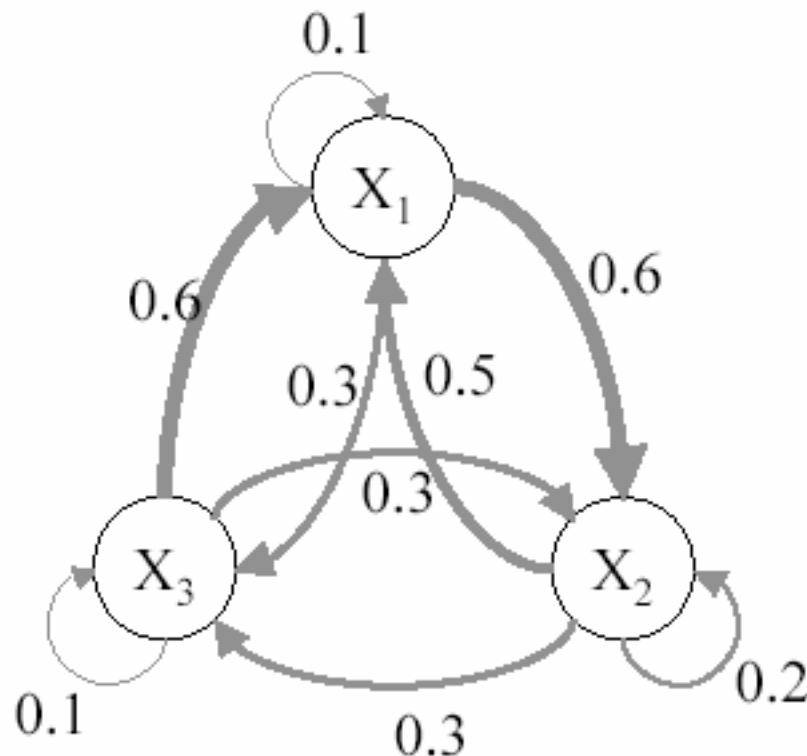
Markov Chain:

- A sequence of random variables X_1, X_2, X_3, \dots
- Each variable is a distribution over a set of states (a,b,c...)
- Transition probability of going to next state only depends on the current state. e.g. $P(X_{n+1} = a | X_n = b)$



transition probs can be arranged in an NxN table of elements
 $k_{ij} = P(X_{n+1}=j | X_n = i)$
where the rows sum to one

A simple Markov chain

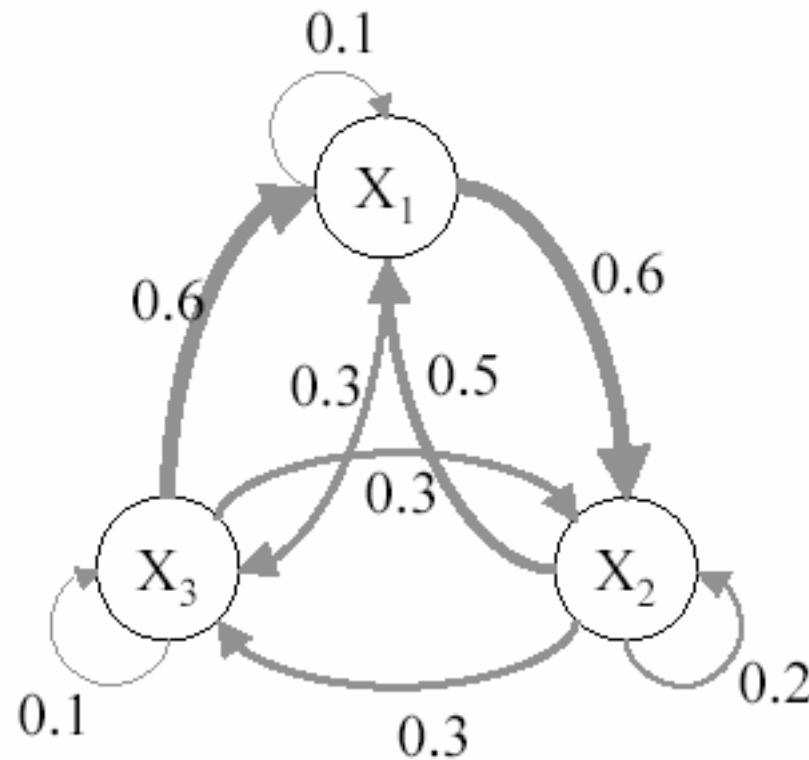


$$K = \begin{bmatrix} & & \\ 0.1 & 0.5 & 0.6 \\ 0.6 & 0.2 & 0.3 \\ 0.3 & 0.3 & 0.1 \\ & & \end{bmatrix}$$

$K =$ transpose of transition prob table $\{k_{ij}\}$ (columns sum to one). We do this for computational convenience.

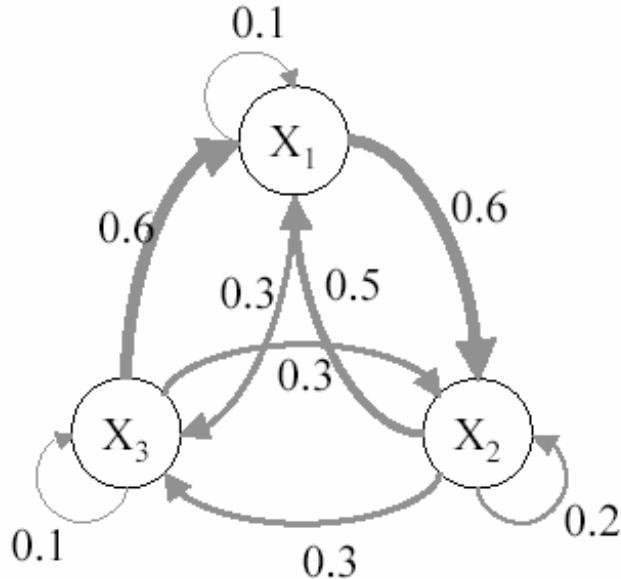
Question:

Assume you start in some state, and then run the simulation for a large number of time steps. What percentage of time do you spend at X_1 , X_2 and X_3 ?



$$K = \begin{bmatrix} 0.1 & 0.5 & 0.6 \\ 0.6 & 0.2 & 0.3 \\ 0.3 & 0.3 & 0.1 \end{bmatrix}$$

Experimental Approach

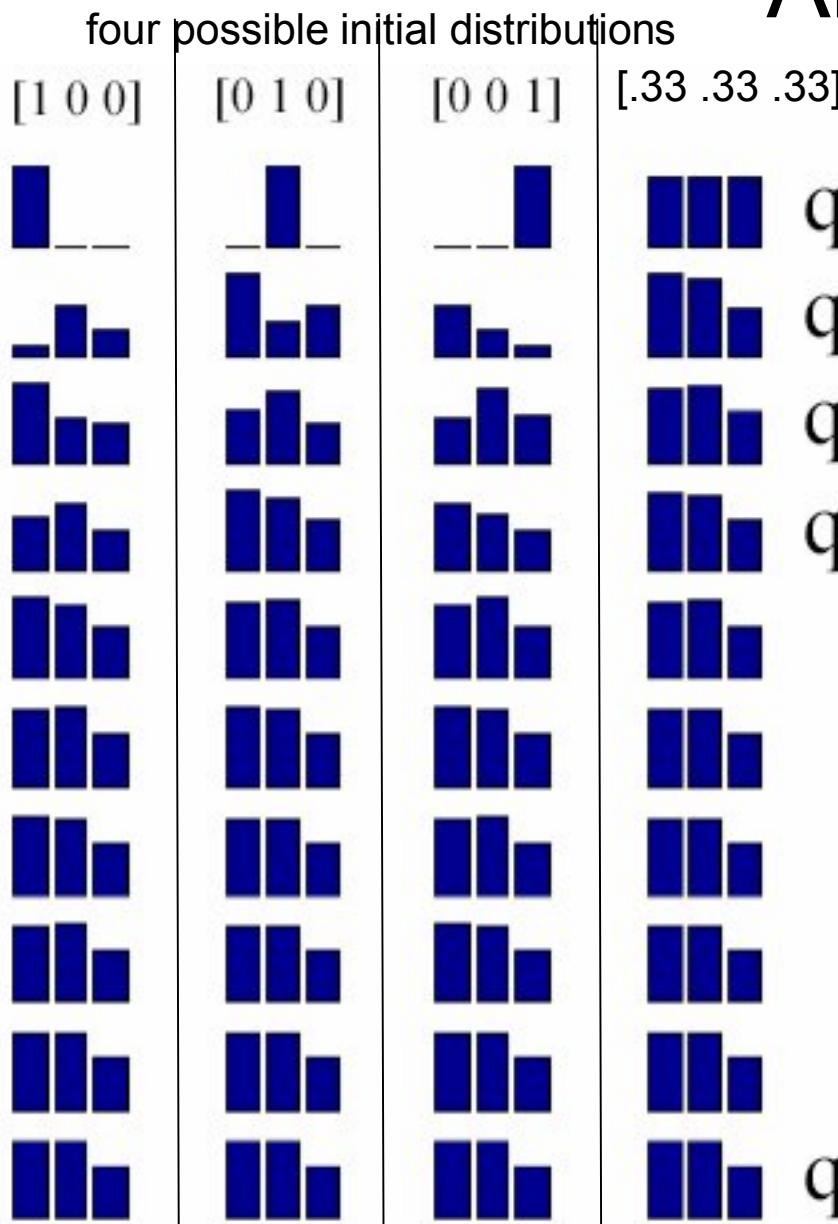


Start in some state, and then run the simulation for some number of time steps. After you have run it “long enough” start keeping track of the states you visit.

{... X1 X2 X1 X3 X3 X2 X1 X2 X1 X1 X3 X3 X2 ...}

These are samples from the distribution you want, so you can now compute any expected values with respect to that distribution empirically.

Analytic Approach



q_0 initial distribution

$q_1 = K q_0$ distribution after one time step

$q_2 = K q_1 = K^2 q_0$

$q_3 = K q_2 = K^2 q_1 = K^3 q_0$

$q_{10} = K q_9 = \dots K^{10} q_0$

all eventually end up with same distribution -- this is the stationary distribution!

Eigen-analysis

K =

0.1000	0.5000	0.6000
0.6000	0.2000	0.3000
0.3000	0.3000	0.1000

$$KE = ED$$

in matlab:
`[E,D] = eigs(K)`

Eigenvalue v_1 always 1

E =

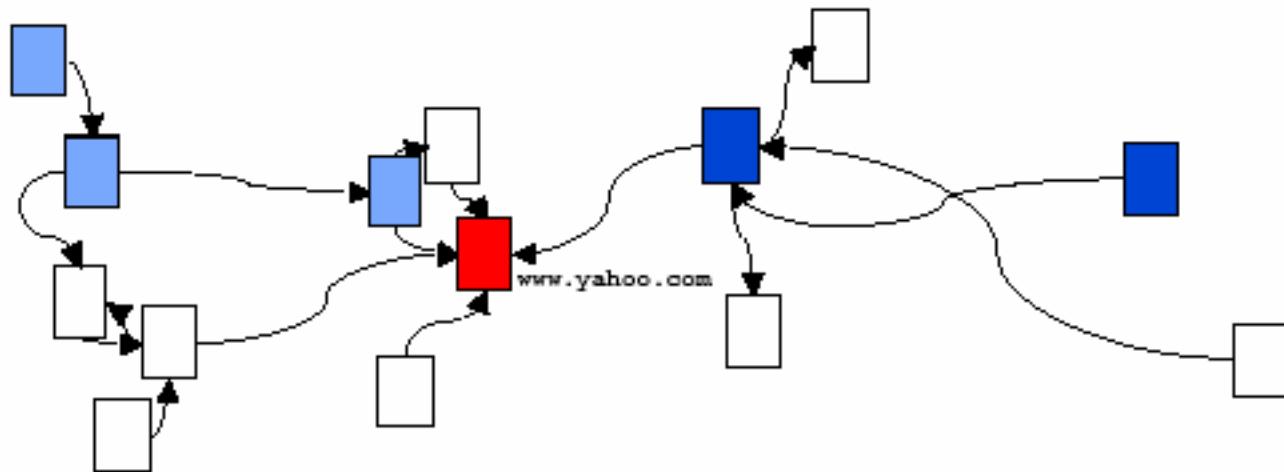
0.6396	0.7071	-0.2673
0.6396	-0.7071	0.8018
0.4264	0.0000	-0.5345

Stationary distribution
 $\pi = e_1 / \text{sum}(e_1)$
i.e. $K\pi = \pi$

D =

1.0000	0	0
0	-0.4000	0
0	0	-0.2000

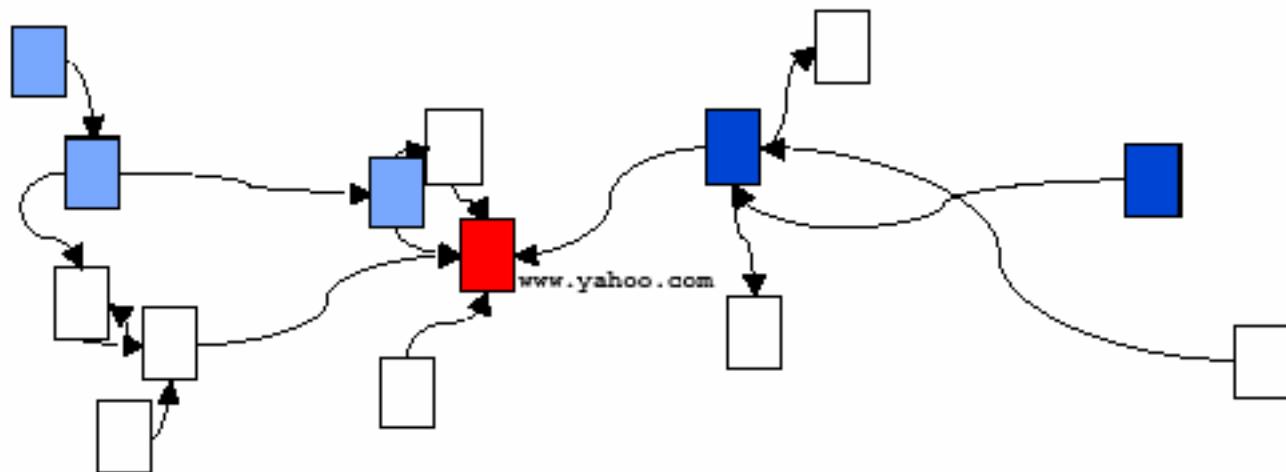
The Web as a Markov Chain



The PageRank of a webpage as used by Google is defined by a Markov chain. It is the probability to be at page i in the stationary distribution on the following Markov chain on all (known) webpages. If N is the number of known webpages, and a page i has k_i links then it has transition probability $(1-q)/k_i + q/N$ for all pages that are linked to and q/N for all pages that are not linked to. The parameter q is taken to be about 0.15.

Google Pagerank

Pagerank == First Eigenvector of the Web Graph !



Computation assumes a 15% "random restart" probability

Sergey Brin and Lawrence Page , The anatomy of a large-scale hypertextual {Web} search engine, Computer Networks and ISDN Systems, 1998

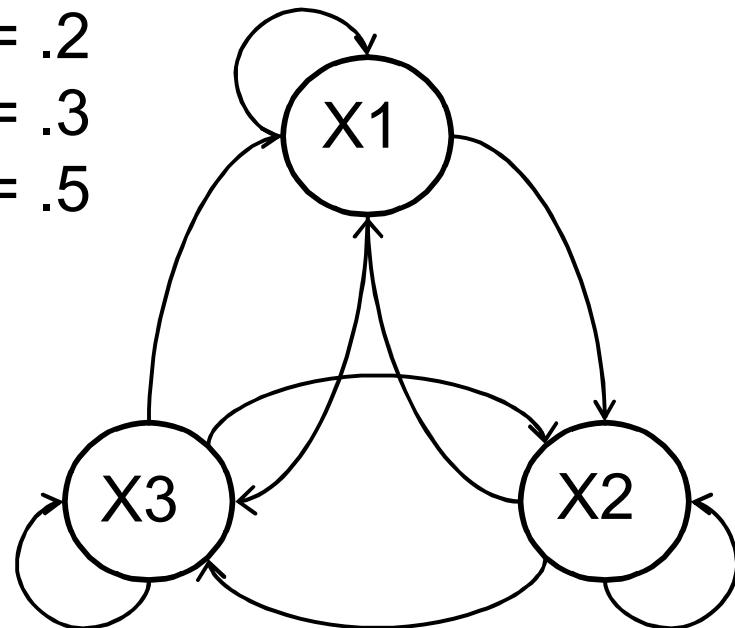
How to Design a Chain?

Assume you want to spend a particular percentage of time at X_1 , X_2 and X_3 . What should the transition probabilities be?

$$P(x_1) = .2$$

$$P(x_2) = .3$$

$$P(x_3) = .5$$



$$K = [\begin{matrix} ? & ? & ? \\ ? & ? & ? \\ ? & ? & ? \end{matrix}]$$

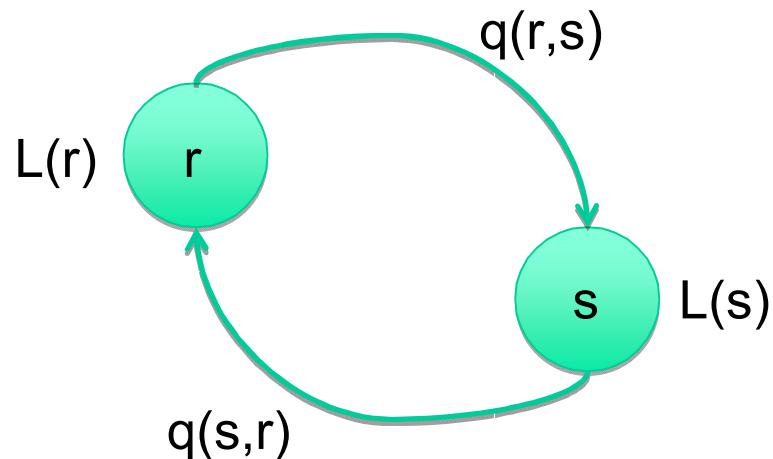
Detailed Balance

- Consider a pair of configuration nodes r, s
- Want to generate them with frequency relative to their likelihoods $L(r)$ and $L(s)$
- Let $q(r,s)$ be relative frequency of proposing configuration s when the current state is r (and vice versa)

A sufficient condition to generate r, s with the desired frequency is

$$L(r) q(r,s) = L(s) q(s,r)$$

“detailed balance”



Detailed Balance

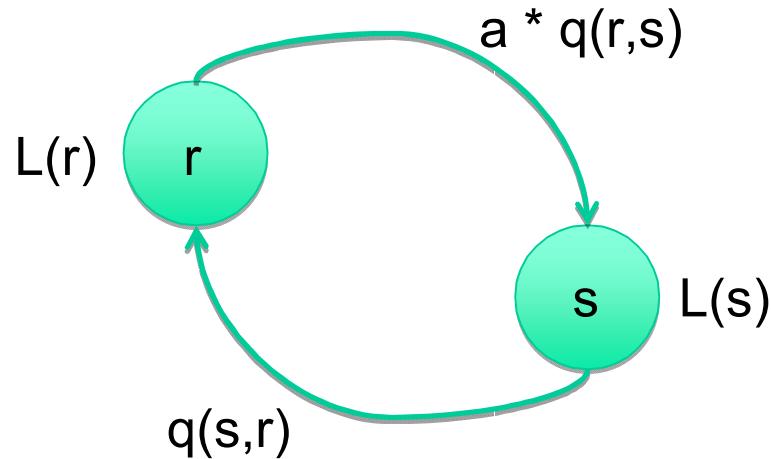
- Typically, your proposal frequencies do NOT satisfy detailed balance (unless you are extremely lucky).
- To “fix this”, we introduce a computational fudge factor a

Detailed balance:

$$a^* L(r) q(r,s) = L(s) q(s,r)$$

Solve for a :

$$a = \frac{L(s) q(s,r)}{L(r) q(r,s)}$$



MCMC Sampling

- Metropolis Hastings algorithm

Propose a new configuration

$$\text{Compute } a = \frac{L(\text{proposed})}{L(\text{current})} \frac{q(\text{proposed}, \text{current})}{q(\text{current}, \text{proposed})}$$

Accept if $a > 1$

Else accept anyways with probability a

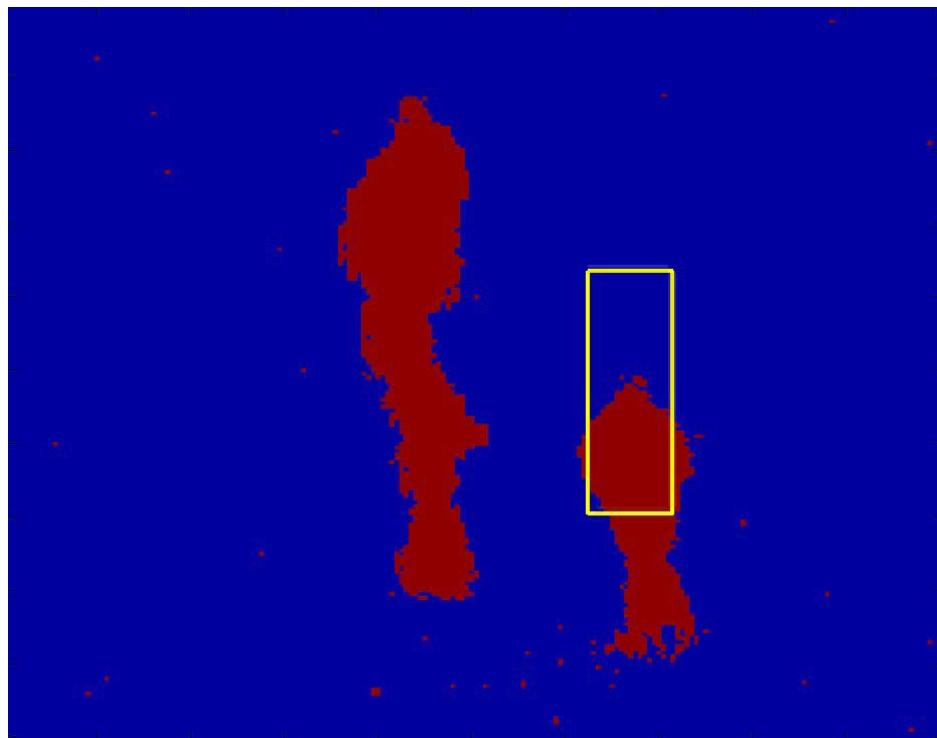
Difference from
Naïve algorithm

Trans-dimensional MCMC

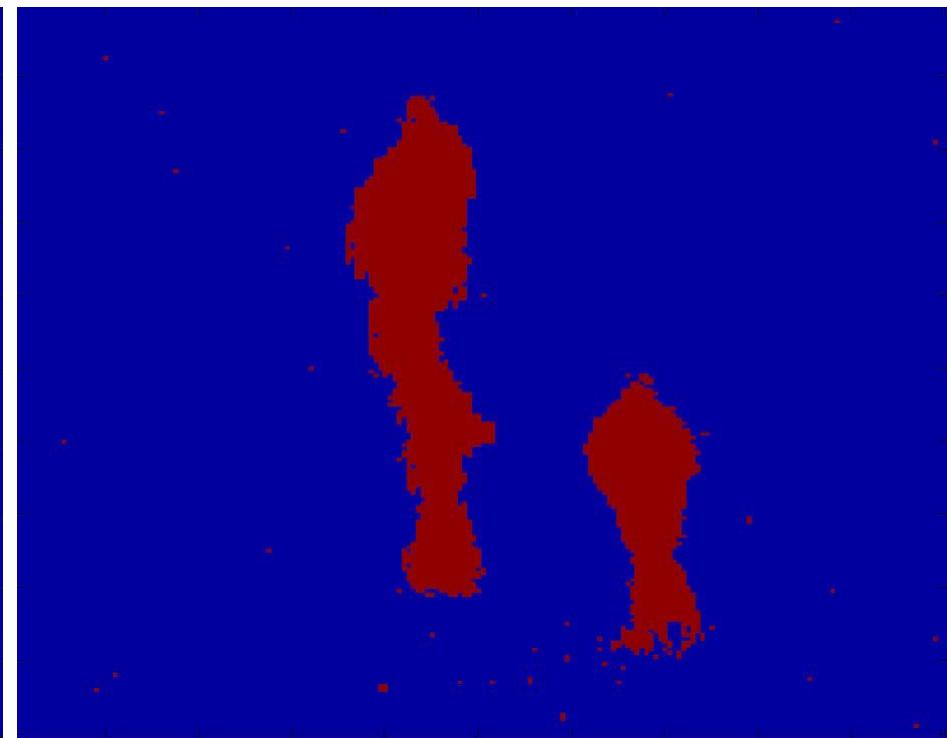
- Green's reversible-jump approach (RJMCMC) gives a general template for exploring and comparing states of differing dimension (diff numbers of rectangles in our case).
- Proposals come in reversible pairs: birth/death and move/move.
- We should add another term to the acceptance ratio for pairs that jump across dimensions. However, that term is 1 for our simple proposals.

MCMC in Action

Sequence of proposed configurations



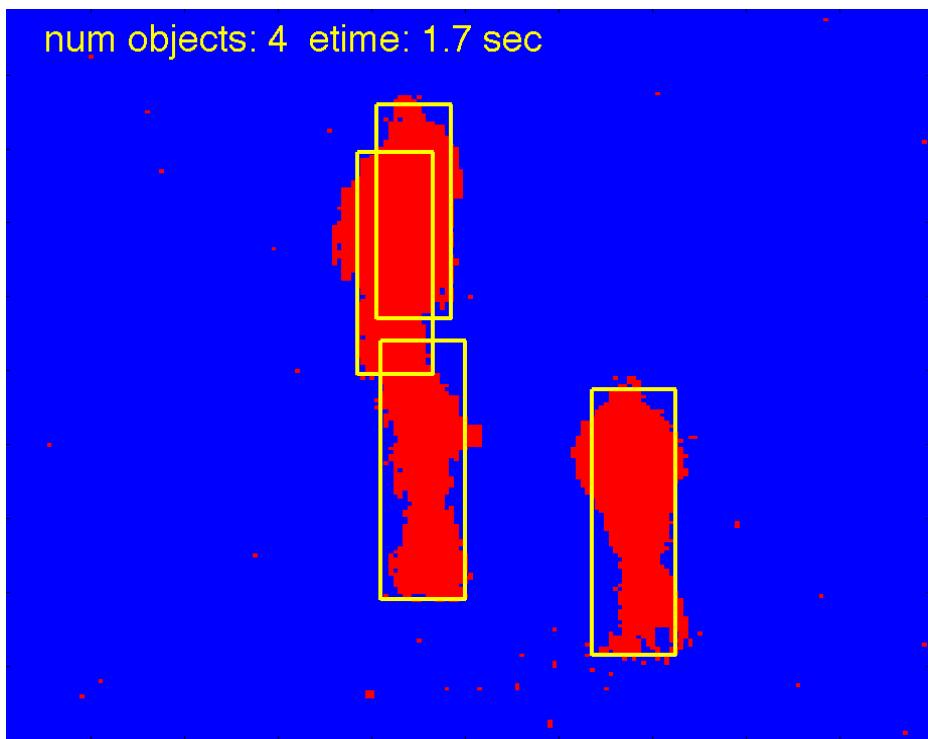
Sequence of “best” configurations



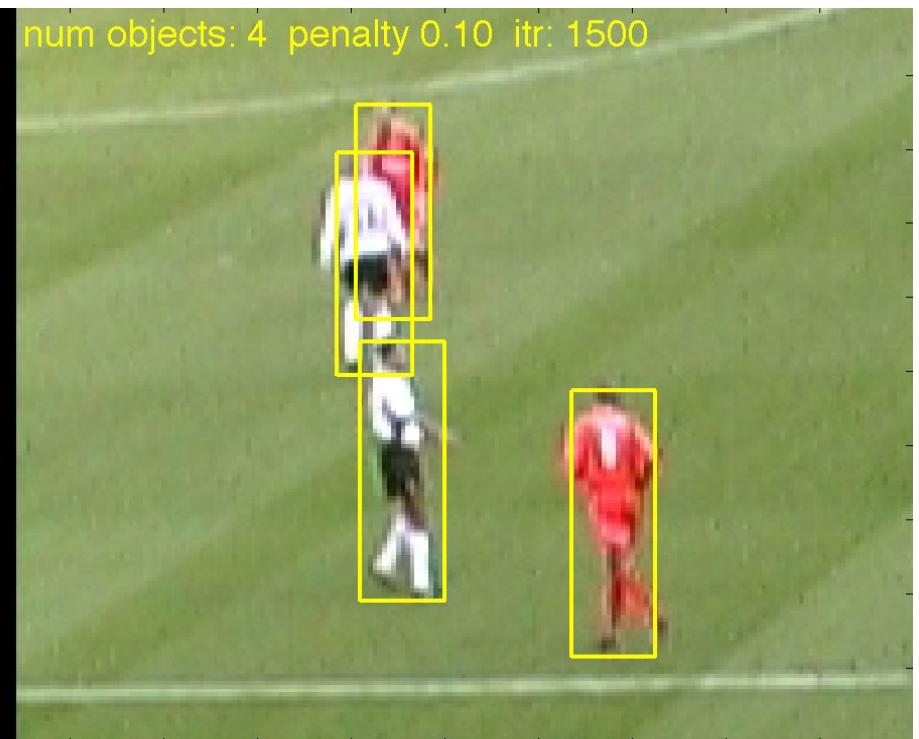
movies

MCMC in Action

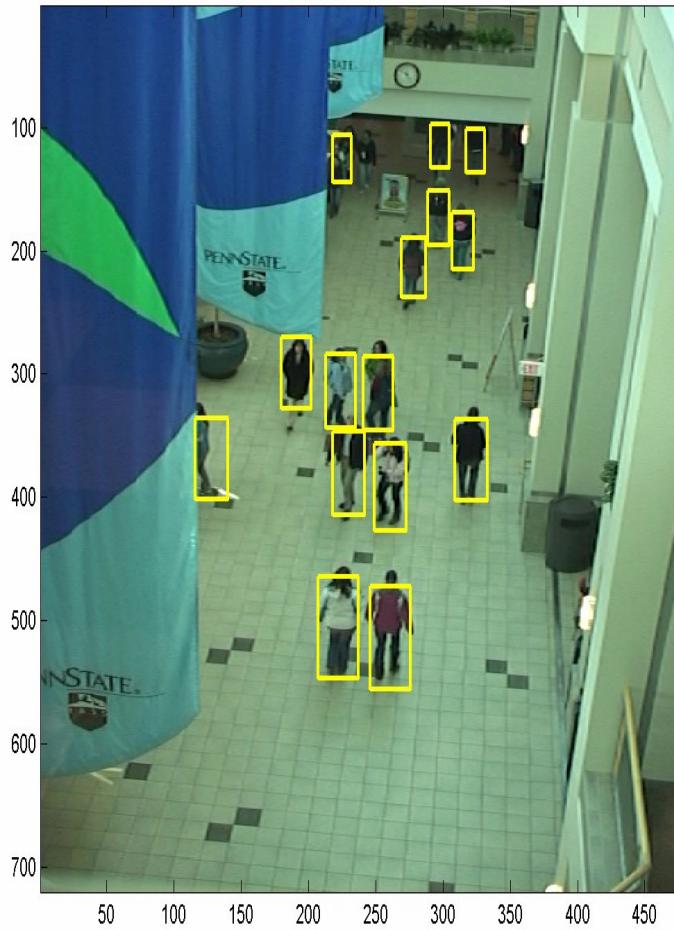
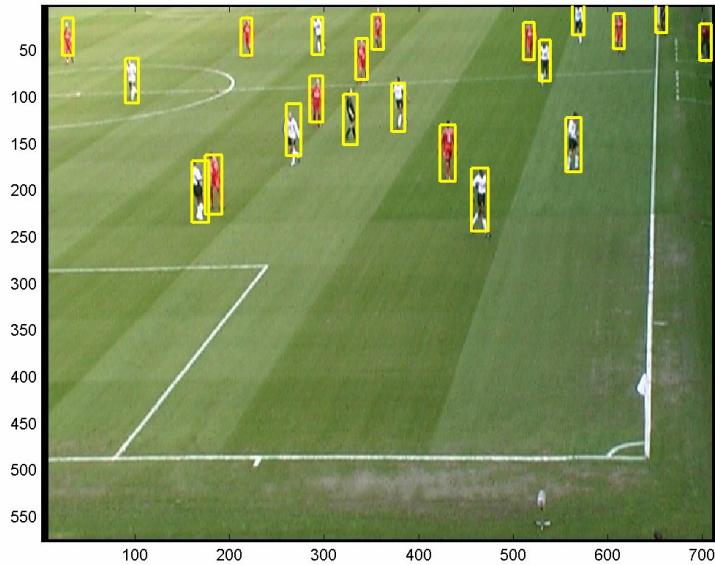
MAP configuration



Looking good!

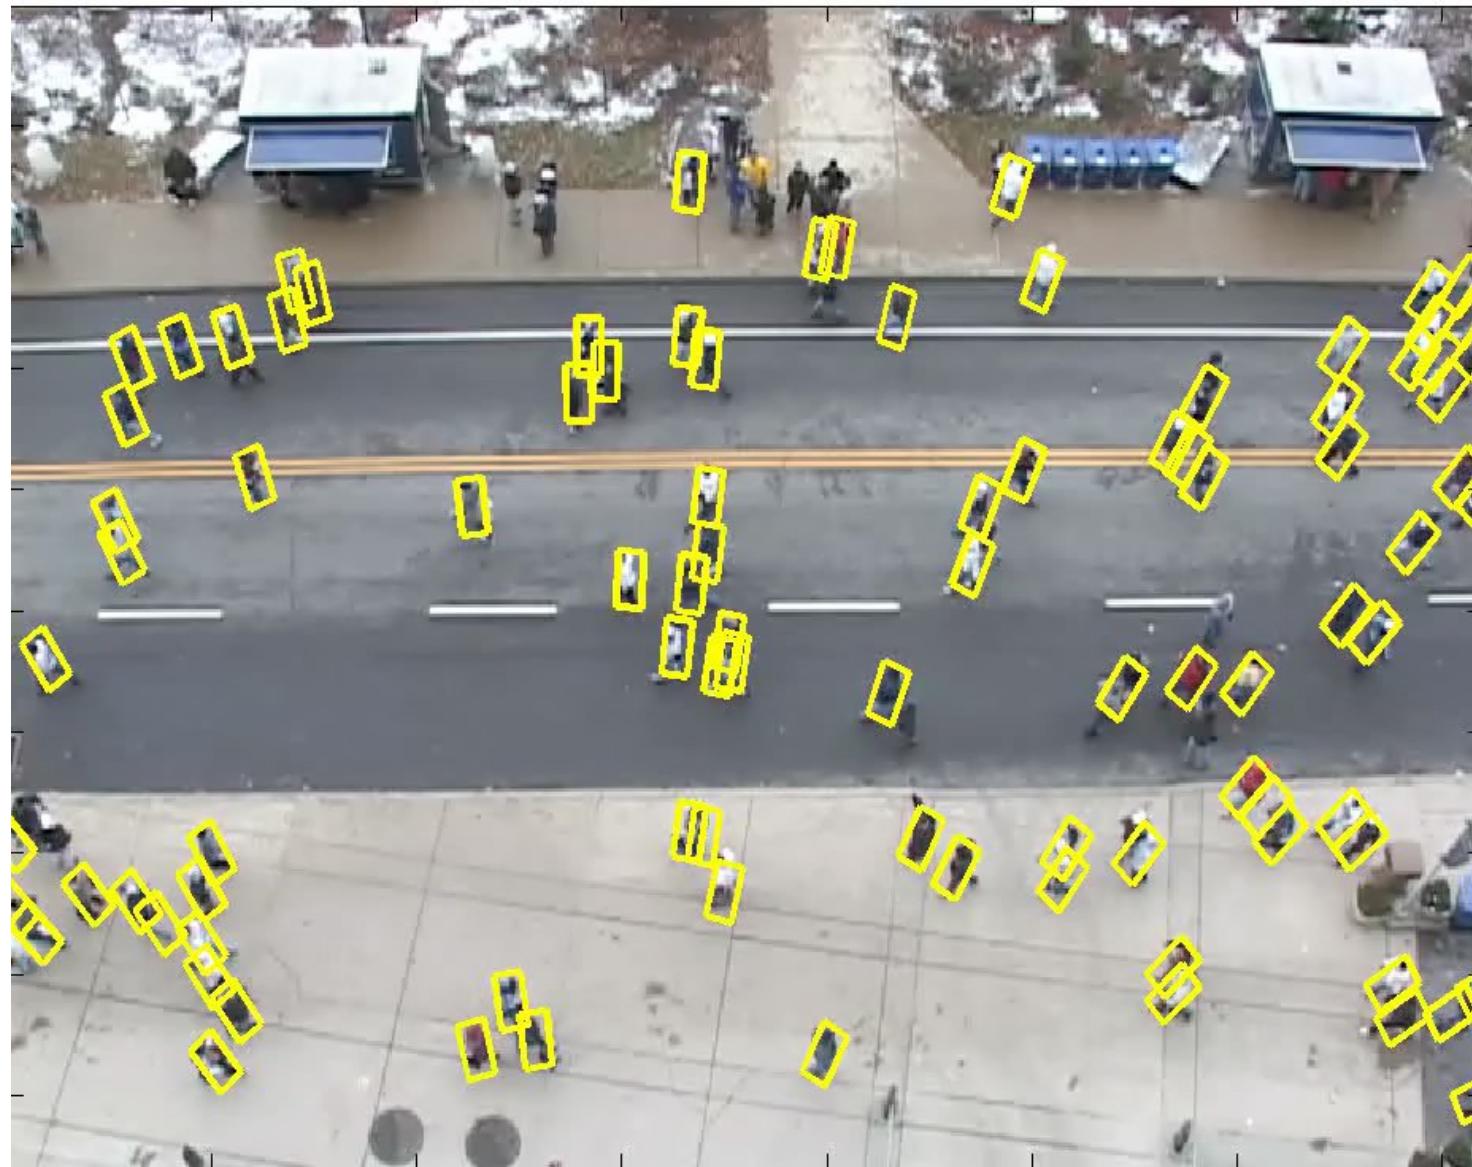


Examples



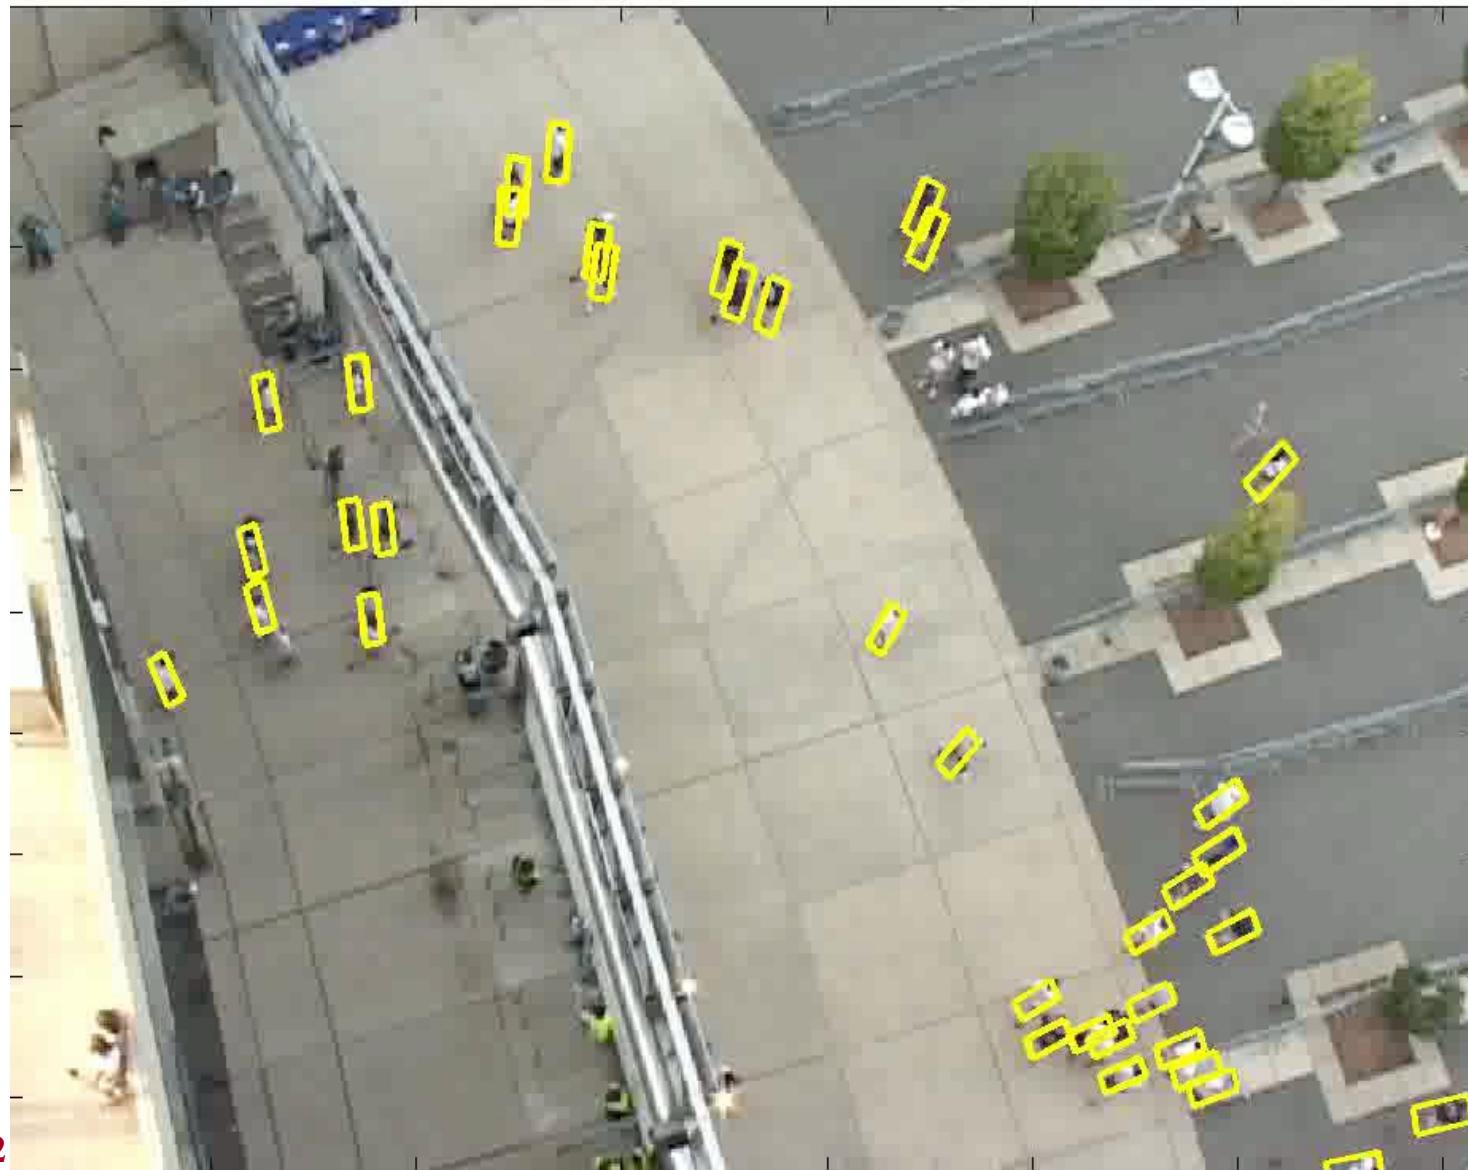
Example: Nov 22, Curtin Road

count 94

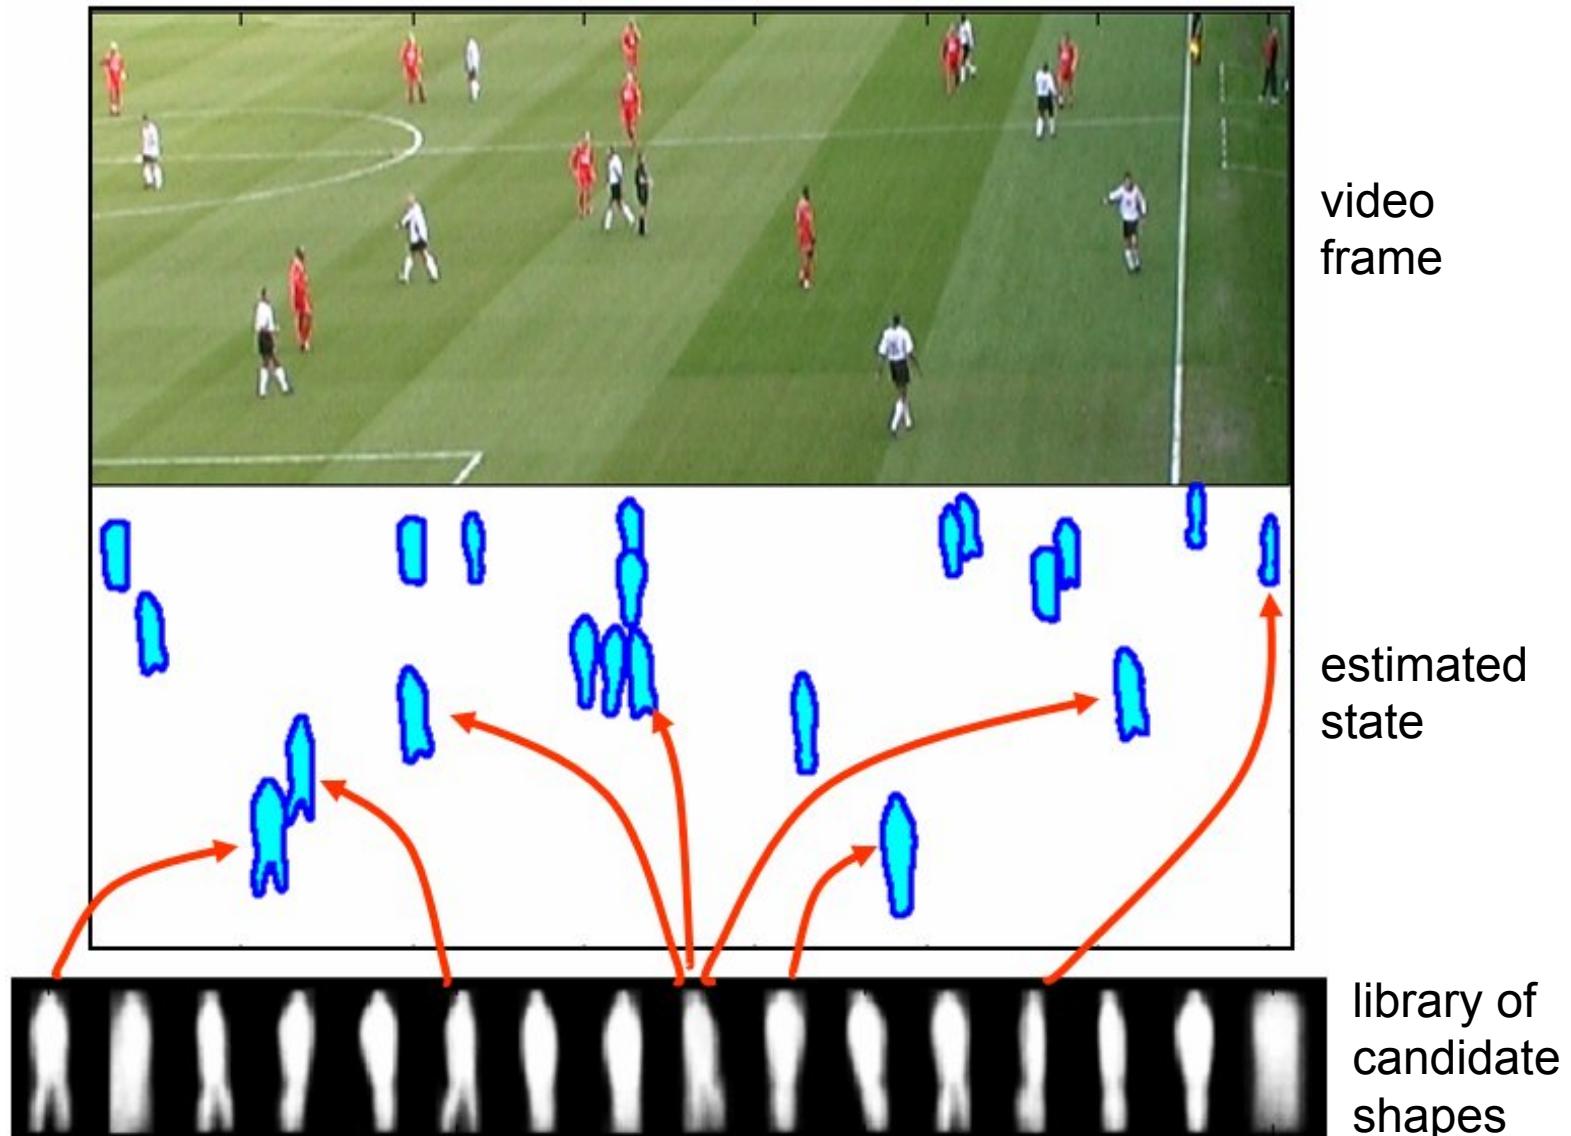


Example: Sep 6, Gate A

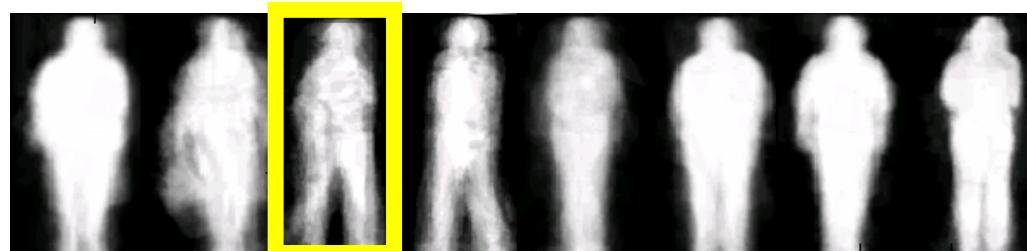
count 37



Adding Shape to the Estimation



Intrinsic vs Extrinsic Shape



Intrinsic shapes:
(silhouettes, aligned
and normalized)

Extrinsic shape
(how bounding box of
silhouette maps into
the image)



Revised Prior Model

$$\pi(\mathbf{o}_i) = \pi(p_i)\pi(w_i, h_i, \theta_i | p_i)\pi(s_i)$$



Prior for
object i



Prior on extrinsic shape
(location + bounding box
height, width and orientation)



Prior on intrinsic
shape selection

Learning Intrinsic Shapes

Silhouette shape represented as a Bernoulli mixture model

$$p(\mathbf{x}|\boldsymbol{\mu}, \boldsymbol{\pi}) = \sum_{k=1}^K \pi_k p(x|\boldsymbol{\mu}_k)$$



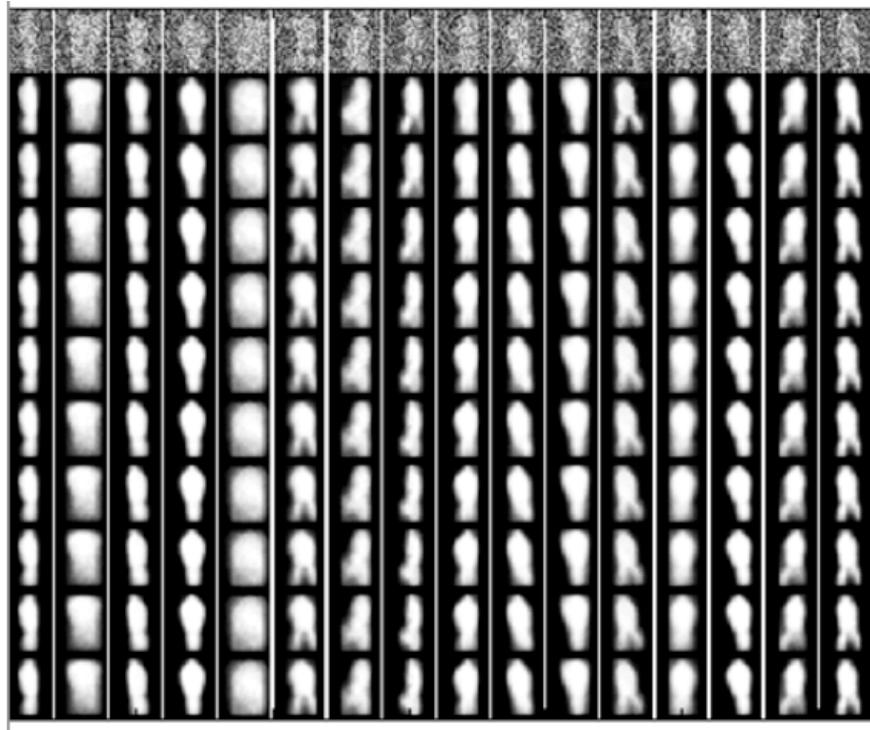
Training shapes
(foreground masks)



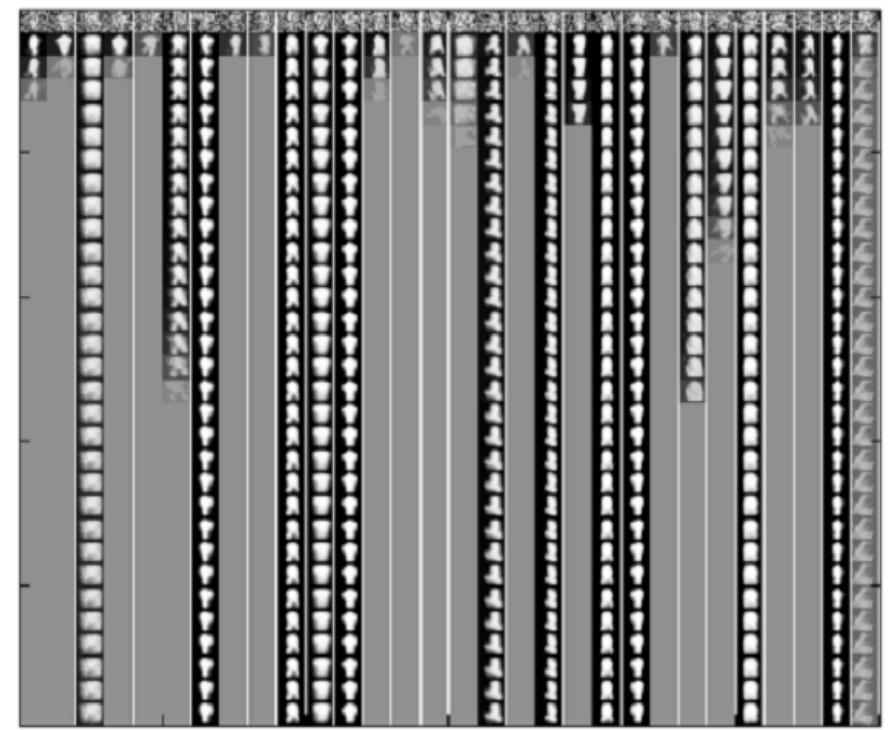
Learned Bernoulli
mixture model
("soft" silhouettes)

Learning Intrinsic Shapes

“standard” EM

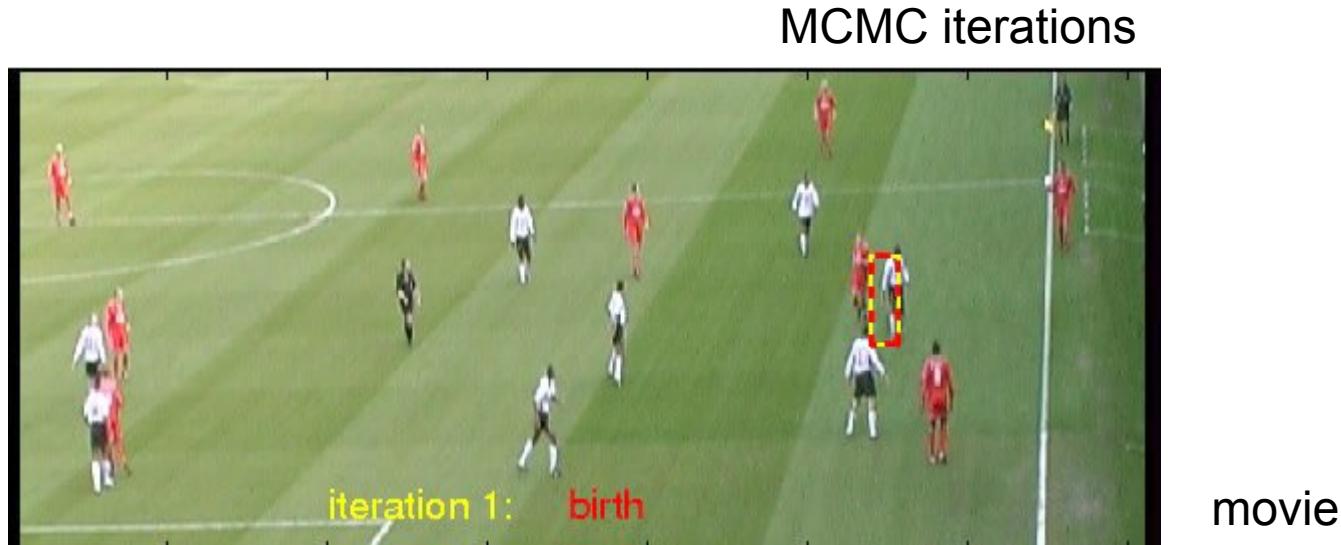


Bayesian EM with Dirichlet prior



Provides automatic selection of number of mixture components

Adding Shape to the Estimation



MCMC proposes changes to current configuration

- add/remove a person
- shift location of person
- change their shape**

Robert Collins

Penn State

VSPETS Soccer Dataset

Evaluation run on every 10th frame
green contours: true positives
red boxes: false negatives
pink contours: false positives



movie

detailed view



V

count=4



count=4



count=4



count=4



count=3



count=4



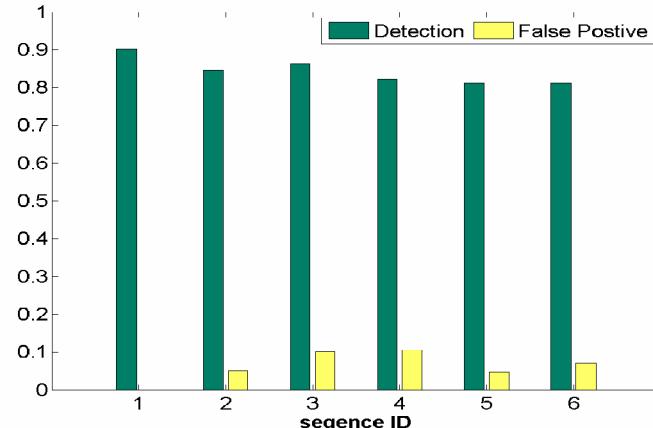
count=4

Robert Collins

Penn State

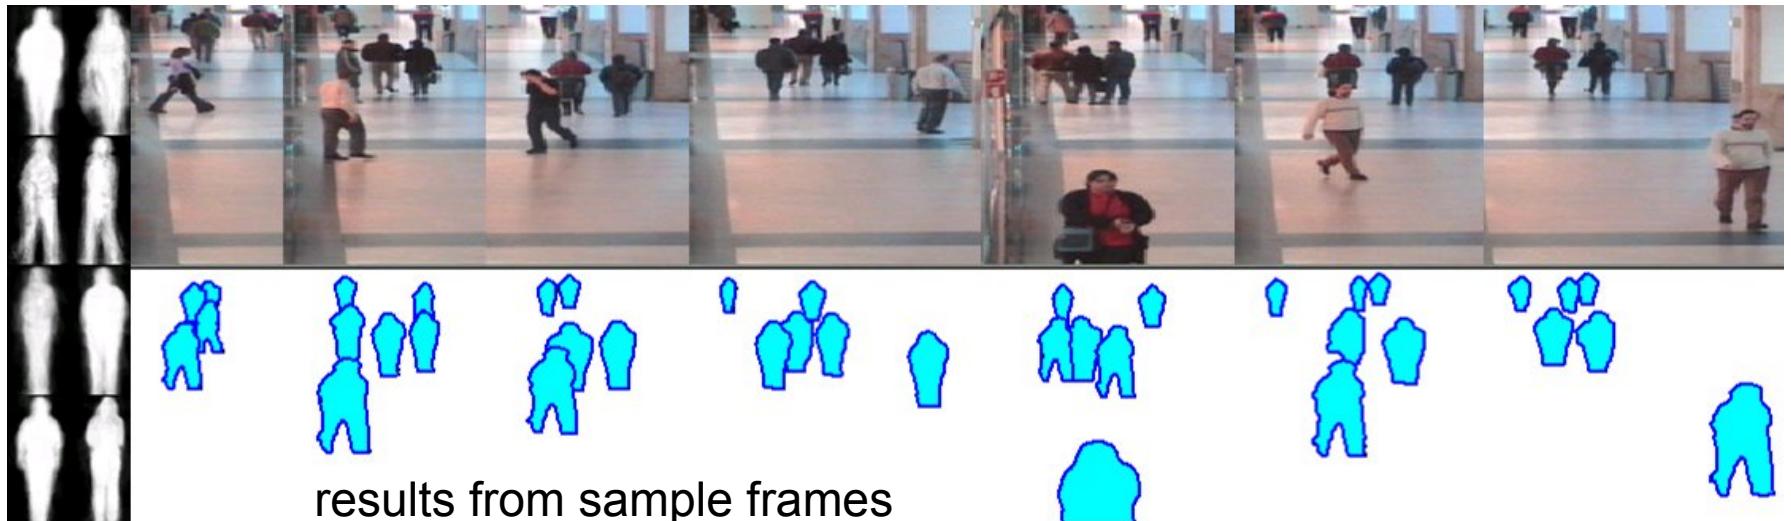
Caviar Dataset

evaluation on six sequences



Dataset	Total # People	Detection Rate	False Positive Rate
CAVIAR	1258	.84	.06
SOCER	3728	.92	.02

Learned shapes



VLPR 2012

CAVIAR sequence #1

average detection rate: 0.90

average false positive rate: 0.00

Contributions

- We introduce a Marked Point Process framework for detecting people in crowds
- Conditional mark process models known correlations between bounding box size/orientation and image location
- Extrinsic and intrinsic shape models are learned automatically from training data
- Bayesian EM is used to automatically determine the number of components in the mixture of Bernoulli model for intrinsic shape

Research Directions

Idea: might be good to try other human shapes (activities)



or other animals/objects.

